Measuring the Task Value of Precision in Closed-Loop Agents

Anonymous ACL submission

Abstract

Evaluations of low-precision agents usually ask how much the model changed. A deployment that keeps a BF16 fallback needs a different quantity: what a specific upgrade is worth under the policy that executes afterwards. An upgrade can replace the current action, the subsequent controller, or both. We present an auditable same-state protocol that measures these interventions separately, counts the rescues and damages each one causes, and compares the BF16 replacement with a same-precision substitute. On Qwen3 agents under simulated quantization in τ^2-bench, replacing a quantized agent's first divergent action changes many task outcomes in both directions with a small net effect on the development panel, and the changes recur on the same tasks under other nonces (exploratory). On held-out telecom tasks under a pre-registered protocol, the net effect is -4.6 points (crossed 90% interval $[-7.2, -2.0]$); the pre-registered readings had no negative category and call this inconclusive. Which upgrade is evaluated changes which rescues are counted: testing only the action misses rescues that need the BF16 continuation as well, and testing only a whole takeover misses action-level rescues that hold under the quantized continuation. A second checkpoint at the same bit width produces comparable turnover (equivalence not shown), so outcome sensitivity is not by itself evidence that more bits are needed, and when all actions are compared at the same states BF16's extra action value is unresolved. Signals of how much the action changes, such as token KL, track whether the action differs rather than what the difference does to the outcome, and a frozen uncertainty trigger enriches outcome-changing decisions without a reliable net gain.

1 Introduction

A quantized LLM agent is cheaper to serve than its BF16 reference, and a deployment can keep the reference as a fallback, as dynamic precision routing does (Li et al., 2026c). Suppose the quantized agent (Q) reaches a decision at which the BF16 agent (B) would act differently. There are three ways to spend B here: **replace only this step, let B take over from the next step, or both**. Average task success does not answer this. On our checkpoints, quantized-minus-BF16 success ranges from -0.2 to $+4.2$ pp across four W4 cells (positive: Q succeeds more often), and equivalence within ± 3 pp is not established in any cell (Appendix A.1); even a precise end-to-end number would not say which step an upgrade should touch.

Figure 1a shows one such decision from our data. In an airline task, Q looks up a reservation id that does not exist and receives "not found"; B looks up the user's profile. We froze the state and ran all four combinations of {B's or Q's action} $\times$ {B or Q continuing}. Only B's action followed by B's continuation succeeds. When B supplies only this step, Q continues, searches for flights, writes a change and fails; when Q's failed lookup stands and B takes over, B fails too. Here the task value of B's action depends on the execution that follows it, and judging the action alone, however carefully, cannot reveal it.

The quantity a fallback decision needs is therefore not how much the model changed, but what a specific intervention is worth under the policy that executes afterwards. Recent work forks live agent trajectories and continues them with a different model, against same-model control forks (Gonuguntla, 2026); there, the new model writes the next action and everything after it, so the two effects cannot be told apart. We cross the source of the current action with the model that executes the rest of the episode, at a frozen state that replays bit-exactly, including the environment databases, the user simulator's view and the remaining step and error budget (§3). A same-precision substitu-

tion, a second W4 checkpoint of the same model in place of B, asks whether the consequences require a higher bit width. We study Qwen3-4B and Qwen3-8B (Qwen Team, 2025) under simulated weight and KV-cache quantization in τ^2-bench (Yao et al., 2024; Barres et al., 2025), BFCL multi-turn (Patil et al., 2025) and ALFWorld (Shridhar et al., 2021), and Mistral-Small-3.2-24B (Mistral AI, 2025) on ordinary BF16-sourced τ^2-bench states. Three distinctions organise the results: an action change does not imply an outcome change; outcome changes are both rescues and damages; and consequences of B's action do not imply more value than a same-precision substitute would bring.

Findings.

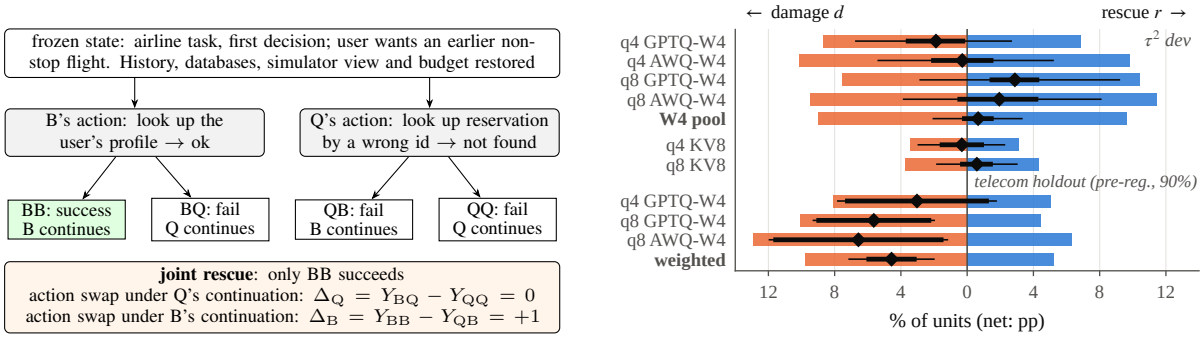
- **A specified upgrade has two-sided consequences and a small net** (§4). Replacing a W4 agent's first tool-structure-divergent action with B's changes the outcome in 15.5–20.9% of task×block units, with rescues and damages of similar size and a net effect of -1.9 to $+2.9$ pp (Figure 1b). On held-out telecom tasks under a pre-registered protocol the net is -4.6 pp, inconclusive by the pre-registered readings. The consequences recur on the same tasks under other nonces (exploratory, post hoc). Action-only and whole-takeover evaluation each miss beneficial interventions that the other counts (Figure 1a), so the two cannot share one gain label.
- **Outcome sensitivity is not evidence that more bits are needed** (§5). A second checkpoint at the same bit width, at W4 and at W3, produces two-sided flips and continuation dependence at comparable size (equivalence not shown), and on common states BF16's additional action value over a second W4 checkpoint is not resolved.
- **Current sensitivity signals do not yet select beneficial upgrades** (§6). Approximate token KL tracks whether the action changes, not what the change does; a frozen Q-only uncertainty trigger captures more outcome-changing decisions than random selection at two of three escalation rates (held-out blocks of the same tasks), but rescues and damages alike, and recovers only 2–5 of the 34–38 net rescues that a hindsight oracle with the same budget finds.
- **Duration and scope** (§7). A takeover's value depends on its duration: under SpinQuant-W4A4, one BF16 decision before handback recovers none of sustained takeover's $+11.8$ pp (exploratory). A key-order experiment on one representation change (300 states) is pending. At ordinary states, on BFCL and ALFWorld, and for Mistral-Small-3.2-24B, the replacement also both rescues and damages.

Contributions. (1) *An executable, auditable same-state intervention protocol* that measures action-only, continuation-only and joint replacement separately, with a same-precision reference substitution. The factorial design is standard; the contribution is making these counterfactuals executable and auditable in closed-loop agents (exact state restoration, budget accounting, replay acceptance, explicit denominators). (2) *Empirical findings*: the value of an upgrade depends on which intervention is evaluated; high turnover hides two-sided consequences and is not by itself evidence that more bits are needed; and current sensitivity signals have not established reliable beneficial selection.

2 Related Work

Quantized agents compared episode by episode. *Flat Score, Amplified Failures* (Jang et al., 2026) is the closest study in setting: on τ^2-bench retail and telecom, 4-bit agents keep a flat task score while process damage grows, with a two-sided paired task split and recovery equally frequent at both precisions. Its episodes are sampled independently, so the precision of the action that went wrong and that of the controller that must recover change together; we hold the state fixed and vary the two separately. Two-sided flips under compression were first shown on static tasks (Dutta et al., 2024); agent-compression benchmarks report averages (Dong et al., 2025).

Forks and counterfactual replay. *The Replay Gap* (Gonuguntla, 2026) evaluates per-step model switching for routers: it forks live SWE-bench Verified trajectories, continues each fork with another model, and compares against same-model control forks. Live forking and a same-model control are therefore not new. Its swap replaces the next action and the continuation together, which is our "both" contrast, and its control holds the model fixed to measure sampling and replay noise. We cross the action source with the continuation controller at a bit-exact frozen state, and our ref-



(a) (b)

Figure 1: Four-branch task value. (a) A real decision (Qwen3-4B, AWQ-W4; airline task 31, decision 1 of 12): B's action is worth a success only if B also continues. This is a *joint rescue* (exploratory: 4.0% of W4 divergent states; the converse pattern, rescue undone, 4.8%; neither differs from a label-permutation reference under the new-task interval, §4). (b) Replacing the first tool-structure-divergent action with B's, per cell: rescues r (right) and damages d (left) in % of all task×block units, and the net $G_{\text{all}} = r - d$ (diamond, pp). Thick/thin: fixed-pool/new-task 95% intervals on the τ^2-bench development panel; block-t/crossed 90% on the pre-registered telecom holdout ("weighted" is its decisional endpoint). All cells: Appendix Table 4.

erence changes the model while holding the bit width fixed. Restoring a state, replacing one action and letting the original policy continue locates decisive errors (Zhang et al., 2025), attributes failures through a do-intervention on one step (Shah, 2026), and turns alternative actions into training signal (Li et al., 2026a,b). This is our "this step only" contrast, and the factorial design itself is standard. We found no work among the 32 closest that crosses the action with the continuation at a precision-induced divergence or reads the result against a same-precision alternative.

Precision routing and its signals. DMR (Li et al., 2026c) learns a per-step BF16/W4 router from BF16-versus-quantized KL supervision and episode reward. Sensitivity has also been localised to phases (Lu et al., 2026) and task types (Zhang et al., 2026), and for quantized reasoning models to high-entropy, high-KL token positions (Lotfi et al., 2026). We do not propose a router. We measure the task value a router should learn, and our simple frozen trigger (§6) does not test, let alone refute, a learned router such as DMR.

Evaluation noise and user simulators. Prefix caching changed multi-turn tool-use trajectories in 36.2% of episodes at 16-bit and 75.0% at 4-bit (Patodiya, 2026); agentic evaluations vary from run to run (Bjarnason et al., 2026; Chen et al., 2026) and with the user model (Seshadri et al., 2026; Jain and Sandhu, 2026).

3 Measuring the Task Value of a Replacement

3.1 Setting

Tasks and agents. We use τ^2-bench in text mode: a development panel of 58 retail and 27 airline tasks plus 80 telecom development tasks, and a holdout of 34 telecom tasks that no agent ran before its protocol was frozen. The agents are Qwen3-4B-Instruct (q4) and Qwen3-8B with thinking disabled (q8); §7 adds Mistral-Small-3.2-24B. One greedy Qwen3-30B-A3B-Instruct user simulator serves all agents. Retail and airline success is the final-database match with normal termination, which does not include all of τ^2-bench's official communication requirements; telecom uses the official task-basis success (Appendix A). BFCL multi-turn and ALFWorld, which have no user simulator, run once per item (Appendix A).

Quantization and serving. Q is a simulated (fake-quantized) checkpoint served with BF16 compute on the same path as B: GPTQ (Frantar et al., 2023) and AWQ (Lin et al., 2024) at W4 and a KV8 cache in the main cells; GPTQ-W3, RTN-W3 and a rotated KV4 cache as boundary (stress) cells. B and Q run on vLLM (Kwon et al., 2023) with batch-invariant kernels (He and Thinking Machines Lab, 2025) and deterministic compilation in separate server processes, so no prefix cache is shared. A *block* is one closed-loop run of Q over the task pool under one prompt nonce (a random session-identifier line in the system prompt); blocks are the replicate unit.

3.2 Four branches on one state

Frozen state. Before Q's decision t we restore the state s_t: the full message history; the environment and user-device databases, rebuilt by strict replay of every recorded tool call; the user simulator's visible conversation; and the used step and error budget. B and Q generate their actions from identical prompt token ids, and the state is written to disk before any branch runs.

Crossed intervention. $Y_{\mathrm{XY}} \in \{0, 1\}$ is task success when the action at s_t comes from X and model Y executes the rest of the episode, X, Y $\in \{\mathrm{B}, \mathrm{Q}\}$ (Figure 1a). The action is executed by the real environment or answered by the real simulator, and all four branches start with the same remaining budget. Each upgrade option is a paired contrast with QQ:

$$
\begin{aligned}
G &= Y_{\mathrm{BQ}} - Y_{\mathrm{QQ}} \quad \text{(this step only)},\\
H &= Y_{\mathrm{QB}} - Y_{\mathrm{QQ}} \quad \text{(from the next step)},\\
F &= Y_{\mathrm{BB}} - Y_{\mathrm{QQ}} \quad \text{(both)}.
\end{aligned}
$$

The paired effect of the action swap under each continuation is $\Delta_{\mathrm{Q}} = Y_{\mathrm{BQ}} - Y_{\mathrm{QQ}}$ $(= G)$ and $\Delta_{\mathrm{B}} = Y_{\mathrm{BB}} - Y_{\mathrm{QB}}$; their difference $I = \Delta_{\mathrm{Q}} - \Delta_{\mathrm{B}}$ is the action–continuation interaction. Δ_{Q} and Δ_{B} take values in $\{-1, 0, +1\}$, so $I \in \{-2, \ldots, +2\}$. When $\Delta_{\mathrm{Q}} \neq \Delta_{\mathrm{B}}$ ("off-diagonal"), the swap's value depends on the continuation; a *strict reversal* is $\Delta_{\mathrm{Q}} = -\Delta_{\mathrm{B}} \neq 0$. For the action swap, a *rescue* is $(Y_{\mathrm{BQ}}, Y_{\mathrm{QQ}}) = (1, 0)$ and a *damage* is $(0, 1)$; with rates r and d, the turnover is $M = r + d$ and the net gain $r - d$. All contrasts measure effects on task success, not which action is correct.

Same-precision alternative. A large M or a frequent interaction could mean that this decision needs higher precision, or only that it is sensitive to any change of model. To tell these apart, we rerun the design with Q′, a second GPTQ-W4 checkpoint of the same model calibrated on a different sample, in place of B: the same Q episodes, tasks and nonce blocks, and a protocol frozen before any outcome (q4 and q8 GPTQ-W4; GPTQ-W3 as a stress check). Turnover M', rescues r', damages d' and the off-diagonal share are defined as above with Q′ for B.

3.3 Making the counterfactuals auditable

Replay acceptance. A causal reading requires exact replay. Before any development outcome was analysed, an engineering gate on real divergence states passed every check: Q→Q no-ops matched in 20/20 states and B→B no-ops in 10/10, in every prompt-token hash, message, database hash, termination, reward and budget counter. In production, every fork task re-ran the QQ no-op on an outcome-independent 10% hash sample of its states, and one mismatch would have quarantined the task's whole output; all 924 sampled development states passed, as did all sampled states of the other banks and of the same-precision run (Appendix A). Exactness is a tested property of this serving stack, not a general one (Yuan et al., 2025): greedy decoding on vLLM 0.29.0 with batch-invariant kernels and deterministic compilation, and simulated checkpoints on dense BF16 kernels; untested conditions and the determinism tests are in Appendix A.

Which states. *Bank P*: for each task, block and cell, the first decision at which a_{B} and a_{Q} differ in kind (tool call, text, stop), tool name, call count or order, or normalised arguments; differences in text only are skipped. *Bank U*: one ordinary decision per task, drawn by a hash fixed before generation, in three nonce blocks (U2B: up to six in §5, §6 and §7), from Q's own trajectories (U1) or from trajectories shared by all cells of a model (U2B: BF16; U2Q: Q).

Denominators and inference. We never mix three denominators. (1) *Full pool*: every task×block unit, units without a Bank P state counting as zero; r, d, M and $G_{\mathrm{all}} = r - d$ use it. (2) *Divergent states*: the Bank P states of a cell; G, H, F and the off-diagonal share use it. (3) *Bank U*: one state per task. Within a cell all branches share each state and all tasks of a block share its nonce, so within-cell contrasts are paired and identified; cross-cell comparisons are descriptive, because each cell selects its own divergences. The primary interval treats the task pool as fixed (block bootstrap with at least 10 complete blocks, else t over block means); a crossed task×block bootstrap is the interval for new tasks. Only complete blocks enter.

4 Specifying the Upgrade: Rescues, Damages and Net

Many rescues, many damages, small net. On the τ^2-bench development panel (85 tasks × 20 complete blocks), replacing a W4 agent's first tool-

Tasks		r	d	M	G_{all}
fragile	est.	21.1	17.6	38.7	$+3.4$
(top 20%)	t	[18.6, 23.5]	[14.3, 21.0]	[34.9, 42.6]	$[-1.0, +7.9]$
	crossed	[14.8, 28.1]	[12.4, 23.3]	[32.3, 45.7]	$[-7.1, +13.7]$
rest	est.	6.5	7.2	13.7	-0.7
	t	[5.2, 7.8]	[6.2, 8.3]	[12.5, 14.9]	$[-2.7, +1.2]$
	crossed	[4.4, 9.0]	[4.8, 9.9]	[10.2, 17.5]	$[-4.1, +2.4]$

Table 1: Repeatable task fragility (exploratory, post hoc; four τ^2-bench W4 cells, equal weight). Tasks are ranked by turnover on one hash-defined half of the 20 nonce blocks and measured on the other half (pp of units; 95% intervals, t over held-out blocks and crossed task×block; direction A to B shown). A replication on other nonces of the same tasks, not a prediction for new tasks. Reverse direction: Appendix E.5.

structure-divergent action with B's, while Q continues, changes the task outcome in $M = 15.5$–20.9% of units (Figure 1b). Rescues and damages are of similar size, $r = 6.8$–11.4 and $d = 7.5$–10.1 pp, so the net gain G_{all} lies between -1.9 and $+2.9$ pp. Two fixed-pool intervals exclude zero (negative for q4 GPTQ-W4, positive for q8 GPTQ-W4); no new-task interval does. KV8 cells turn over fewer outcomes ($M = 6.5$–8.1%, a descriptive comparison). G_{all} is a local replacement effect, not an end-to-end comparison (Appendix A.1). Pooled over the four W4 cells, $G_{\text{all}} = +0.7$ pp (fixed pool $[-0.3, +1.6]$; new tasks $[-2.1, +3.5]$) with $M = 18.6\%$. On 34 unseen telecom tasks, under a protocol frozen before any holdout unit ran (12 fresh blocks, three W4 cells), the replacement again rescued and damaged ($r = 5.2$, $d = 9.8$ pp). The primary $G_{\text{all}} = -4.57$ pp (crossed 90% $[-7.17, -1.96]$) is *inconclusive* by the pre-registered readings, which had no negative category; outside them, the interval lies below zero (Appendix F). On the 80 telecom development tasks the same q8 cells give $+2.7$ and -0.6 pp (Appendix E.5), so the holdout differs from the retail/airline panel in both domain and tasks.

The consequences recur on the same tasks. We split the 20 nonce blocks into two halves by an outcome-independent hash of the block name (written before this post hoc analysis loads outcomes), ranked tasks by turnover on one half and measured the other half (Table 1). The 17 most fragile tasks (top 20%) carry 41.0% (crossed [26.6, 55.9]) and 42.3% ([27.0, 57.5]) of the held-out turnover in the two directions, against 20% if fragility did not persist. Their turnover is 38.7%, while the other 68 tasks still turn over 13.7%. Both groups have rescues and damages; neither group's net gain is resolved under the crossed interval (Appendix E.5). This replicates fragility on other nonces of the same tasks; it does not predict which new tasks are fragile.

Which mode is evaluated decides which rescues are counted. On divergent states, the average effect of the action swap has the same point estimate under Q's and B's continuation (W4 pool: $\Delta_{\text{Q}} = G = +0.5$, fixed pool $[-0.7, +1.7]$, new tasks $[-2.9, +3.8]$; $\Delta_{\text{B}} = +0.5$, $[-1.0, +2.1]$, $[-3.5, +4.5]$ pp), and no average gain is detected for upgrading this step, the rest, or both ($H = +0.2$, $F = +0.7$ pp, all intervals covering zero). The averages hide states on which the single-mode evaluations disagree (exploratory, post hoc; % of W4 divergent states; fixed-pool, then new-task 95% intervals). In a *joint rescue*, 4.0% ([3.5, 4.5]; [2.7, 5.3]), only Y_{BB} succeeds (Figure 1a): testing only the current action misses states where B's action needs B's continuation too. In a *rescue undone*, 4.8% ([4.1, 5.4]; [3.2, 6.5]), B's action rescues under Q's continuation ($Y_{\text{BQ}} = 1$, $Y_{\text{QQ}} = 0$) but B's takeover undoes it ($Y_{\text{BB}} = 0$): testing only a whole takeover, which swaps the action and everything after it, misses an action-level rescue that holds under Q's continuation. A conditional label-permutation reference (which keeps how many of the four branches succeed) already produces both patterns at similar size (joint rescue 4.0% vs 3.5%; rescue undone 4.8% vs 5.8%); the observed frequencies give no generalisable evidence of extra enrichment (Appendix B.2). The reference does not model execution, so it does not say how these outcomes arise. The patterns matter as an evaluation consequence: action replacement and whole takeover cannot share one gain label. Paired comparisons of the three modes, by cell, are in Appendix B.1.

5 Does the Evidence Call for More Precision?

A large turnover or a frequent interaction is tempting evidence that a state needs higher precision. If these consequences required a higher bit width, the same-precision substitution (§3.2) should produce markedly fewer of them. Because each design forks at its own first divergence, the two may select different states, so shared and exclusive flips cannot be split into effects caused and not caused

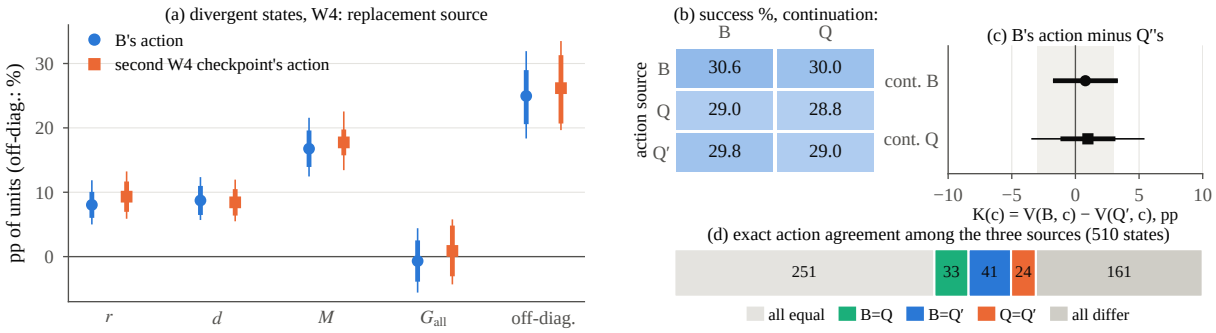


Figure 2: Does an upgrade need a higher bit width? (a) Each design at its own first divergence, W4 (q4 and q8 GPTQ-W4 pooled; 1020 task×block units; $K = 6$ blocks): replacing Q's action with B's (circles) or with a second W4 checkpoint's (squares); rescues r, damages d, turnover M and net G_{all} in pp of units, off-diagonal share in % of divergent states; thick: fixed-pool 95% (t over block means), thin: new-task 95% (crossed). A second W4 checkpoint's action turns over about as many outcomes as B's. (b–d) Common states (q8; 510 pre-registered ordinary states, 6 blocks): (b) success (%) by action source and continuation; (c) $K(c)$, B's action value beyond a second W4 checkpoint's; here, unlike (a), thick: task-cluster bootstrap 95%, thin: t (block means) 95%; shaded: ±3 pp; (d) exact agreement of the three actions. BF16's extra action value is unresolved. Overlap of affected units, equivalence tests, per-model numbers and two-look intervals: Appendix C.

by precision.

A higher bit width is not necessary for two-sided flips and continuation dependence. They also appear when only the calibration checkpoint is swapped (Figure 2a). At W4 (frozen final analysis set, $K = 6$ blocks), the second checkpoint's action rescues $r' = 9.3$ (fixed pool [7.0, 11.7]; new tasks [5.9, 13.2]) and damages $d' = 8.4$ ([6.4, 10.5]; [5.5, 12.0]) pp of units ($M' = 17.7\%$ against $M = 16.8\%$ for B on the same 6 blocks and two GPTQ-W4 cells), and 26.2% of its divergent states are off the diagonal, against 25.0% for B; each same-precision rate has a fixed-pool 95% lower bound above zero. In the W3 stress pair the turnovers are 17.3% (second W3 checkpoint) and 18.8% (B) and the off-diagonal shares 27.6% and 27.3%. Because a same-precision substitute produces turnover and off-diagonal shares of this size, flip rate and interaction size alone cannot diagnose a need for a higher bit width. We do not claim the turnovers are equal: neither pair is equivalent within ±3 pp (Appendix C).

Units affected by both, or by only one, substitution. At W4 (Appendix C), 107 units flip under both substitutions, 64 only under B and 74 only under the second checkpoint. The reference that matters keeps each task's flip propensity: a within-task block shuffle gives 76 [69, 84] units flipping under both, so the substitutions share affected units beyond the fact that the same tasks are flip-prone anyway. The shared core is on the action side. We measure it with two Cohen's κ statistics over the 1020 W4 units: κ_r compares the unit-level indicators "B's action rescues" and "the second checkpoint's action rescues", and κ_d the corresponding damage indicators; $\kappa_r = 0.60$ and $\kappa_d = 0.54$. When B also takes over the continuation, the affected units are largely substitution-specific (Appendix C). The difference in action-side net value (ΔG_{all}) comes only from the substitution-specific units and is not resolved (−1.6 pp; fixed pool [−6.2, +3.0], new tasks [−6.5, +3.6]). These analyses are descriptive (Appendix C).

Common states: sensitivity, extra value and gain are different quantities. A pre-registered comparison holds the states fixed: on 510 ordinary q8 states (85 tasks × 6 blocks), the current action comes from B, from Q or from a second W4 checkpoint Q′, and B or Q continues (Figure 2b–d). Under Q's continuation, replacing Q's action by B's flips 7.1% of outcomes for a net +1.2 pp; replacing it by Q′'s flips 8.0% for a net +0.2 pp. What B's action adds beyond another W4 action, $K(c) = V(\mathrm{B}, c) - V(\mathrm{Q}', c)$, is +1.0 pp under Q's continuation and +0.8 pp under B's. After the fixed extension from three to six blocks, no advantage is resolved (95% t [−3.5, +5.4] and [−1.8, +3.3]). Under B's continuation the ordinary 90% interval lies within ±3 pp (two one-sided tests, t [−1.2, +2.8]), but the two-look sensitivity analysis (94% intervals) does not, so we do not claim equivalence; q4 (two blocks) is also not resolved (Appendix C). A second W4 action changes about

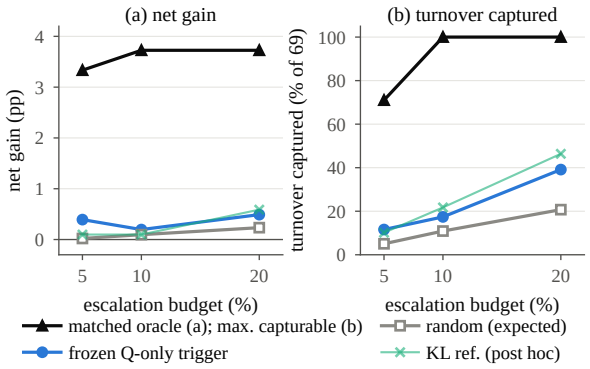


Figure 3: Sensitivity versus benefit, by escalation budget, on held-out ordinary states (1020 U1 rows; 38 rescues, 31 damages): a frozen Q-only uncertainty trigger, matched-count random selection (expected), a hindsight oracle with the same per-cell counts, and a faint post-hoc KL(B‖Q) reference (needs BF16). (a) Net gain, (rescues − damages)/1020, pp. (b) Share of the 69 outcome-changing states selected; black: the most any selection with the same per-cell counts can capture. Values: Table 13.

as many outcomes as B's, and Q′ against Q under Q's continuation is within ±3 pp by the same tests (descriptive). How sensitive the loop is to an action, what higher precision adds at a state and what a given upgrade gains must therefore be measured separately.

6 Can Sensitivity Signals Select Beneficial Upgrades?

§4 and §5 show that outcomes are sensitive to the action and that sensitivity alone does not identify a need for more bits. A router needs more: a signal that picks the upgrades with positive net value. We test a fidelity diagnostic on shared states and a frozen Q-only trigger.

KL tracks whether the action changes, not what the change does. On ordinary BF16-sourced states shared by every method (U2B), we compare methods by $D_{cc} = Y_{\mathrm{QB}} - Y_{\mathrm{BB}}$, Q's action against B's under the shared BF16 continuation (Appendix D). Across the eleven q8 cells (exploratory, descriptive Spearman), mean teacher-forced approximate token KL(B‖Q) follows the action-difference rate ($\rho = 0.99$); its relations with the flip rate given a difference (0.46) and with the net D_{cc} (-0.44) are weak. Its correlations with the rescue rate ($\rho = 0.81$) and damage rate ($\rho = 0.93$) run through the difference rate. Within cells it separates differing from identical actions (mean AUC 0.93) but not flips among them (0.46). SpinQuant-W4A4 is closer to BF16 than QuaRot-W4A4 (KL 0.165 against 0.199), yet the pre-registered contrast of their task value is $+1.2$ pp (t, 3 blocks $[-3.9, +6.2]$; crossed $[-4.7, +7.1]$; Holm-adjusted $p = 1.0$; one block rests on an A7 re-run). An error-dose scan agrees: action agreement with BF16 falls at every dose step, while success responds only at the largest steps of one stress quantizer (Appendix Table 12).

Frozen trigger. G, H and F require running B and all four branches; a deployable rule sees only Q. We test whether a Q-only signal that finds outcome-changing decisions also finds the ones worth upgrading. Scored at the token level on one development block, Q's entropy at the start of its action separates states where B and Q act differently from states where they act identically (AUC 0.66–0.91 per sub-bank and model), but among states where the actions differ, no Q-visible statistic reliably separates flips from non-flips (Appendix D.4). We then froze a Q-only trigger on that block before any later state was scored: escalate one decision when the minimum Q margin over its action is low (margins predict cross-precision divergence; Du et al., 2026), with per-cell thresholds at 5, 10 and 20% escalation rates. On two held-out blocks (1020 rows, 69 outcome-changing states) it captured 8, 12 and 27 outcome-changing states, against 3.48, 7.52 and 14.30 expected at random ($p = .0182, .0613, .0002$; Figure 3b).

It does not yet select net-beneficial replacements. The captured states were rescues and damages alike (6/2, 7/5 and 16/11), for net gains of $+4$, $+2$ and $+5$ against random expectations of $+0.21$, $+0.98$ and $+2.38$ (exact one-sided $p = .0345, .417, .262$); no escalation rate was designated as primary in advance, and the 5% result does not survive Holm correction over the three rates. A hindsight oracle with the same per-cell counts nets $+3.33$, $+3.73$ and $+3.73$ pp (Figure 3a). This headroom is positive by construction, because the oracle picks only rescues and never a damage; the trigger captures 4/34, 2/38 and 5/38 of it. A post-hoc KL(B‖Q) score, which needs BF16, also captures turnover above chance without a significant net gain (Appendix D.5). Enriching outcome-changing decisions is therefore not enough for reliable net gain. The held-out blocks reuse the development tasks, and the test concerns

this simple frozen rule: a learned router trained on task reward, such as DMR (Li et al., 2026c), is not tested here.

7 How Long and Where Do the Consequences Hold?

Takeover duration. On 255 q8 U2B states of the §6 KL panel (3 blocks; exploratory), Q's action is followed by L BF16 decisions and then by Q again. Under SpinQuant-W4A4, BF16's continuation is worth $+11.8$ pp when it runs to the end (crossed 95% [5.1, 18.4]), but one BF16 decision before handback is worth 0.0 [-3.5, $+3.5$], and four decisions $+6.3$ [0.0, 12.2], still 5.5 pp short of sustained takeover [-10.6, -0.8]. Four-decision windows handed control back at 66% of states; the rest ended inside the window, so they cannot show that four steps suffice. Under GPTQ-W4 the continuation value itself is small ($+0.8$ [-5.9, $+7.5$]) and no window effect is resolved (Appendix E.6). The value of a takeover thus depends on how long it lasts, not only on where it starts.

Representation change. *[Pending: key-order reversal; 300 verified tool-call states, 41 tasks.]*

Ordinary states and the user simulator. At ordinary (Bank U) states, turnover is lower (W4 cells, U1 and U2B: 4.7–9.0%), the net gain covers zero, and most outcome changes come from differences in text, whose labels did not reproduce under a second user simulator, so we do not report them as a property of the action (Appendix E.1). On the shared q8 U2B panel used for KL (§6), B's action alone is resolved in no cell, while fixed-pool intervals for B's takeover or full substitution lie above zero in several (Appendix D).

Other environments and a second family. On BFCL multi-turn and ALFWorld, which have no simulated user (one run per item), the replacement turns over 5.2–24.6% of items in both directions (post hoc; Appendix E.4). On Mistral-Small-3.2-24B, at its own ordinary BF16-sourced states (83 tasks, a dated amendment), B's action in GPTQ-W4 rescues (fixed-pool 95%) $r = 4.6$ [2.4, 6.8] and damages $d = 5.4$ [2.4, 8.5] pp ($M = 10.0$ [7.0, 13.1]; RTN-W4 $M = 10.4$), and no mode's net value is resolved ($G = -0.8$ [-5.2, $+3.6$], $H = +1.4$, $F = +0.0$; Appendix E.3). Stress cells are kept separate from W4 (Appendix A.2).

8 Discussion

From model change to intervention value. The results separate three questions that a single success rate or flip rate mixes. How often an upgrade changes the outcome is not how much it gains: B's action changes many outcomes, in both directions, for a small net, and the consequences recur on the same tasks (§4). Which upgrade is evaluated changes which rescues are counted, so action replacement and whole takeover need separate gain labels. And how sensitive the loop is to a change of action is not what higher precision adds: a second checkpoint at the same bit width changes outcomes about as often, and on common states BF16's additional action value is unresolved (§5). A signal of how much the action changes is not a signal of what the change does, and a signal that finds outcome-changing decisions does not yet find net gains (§6). We do not conclude that precision is irrelevant: neither reference establishes equivalent turnover, and the resolved fixed-pool cells and the negative holdout estimate (§4) remain.

Why a fallback cannot be judged by its average gain. A post hoc decomposition of the τ^2-bench W4 pool shows how the small net is composed. Among task×block units where QQ fails (units without a divergent state count as no rescue), replacing the action rescues 14.6%; among units where QQ succeeds, it damages 26.1%. The totals are similar partly because QQ fails more often than it succeeds ($P(\text{QQ success}) = 34.6\%$ of units). An average gain therefore hides a rescue rate on failures and a larger damage rate on successes, and a fallback evaluation has to report both. This decomposition conditions on outcomes known only afterwards, so it cannot be used online to decide when to escalate.

Recommendations for evaluating precision interventions in agents. (i) Specify the intervention (the action, the continuation, or both) and report its rescues, damages and net on the same states, not only success or flip rates. (ii) Include a same-precision reference substitution, so that sensitivity is not read as a need for precision. (iii) Give intervals for both the fixed task pool and new tasks (crossed), because they can differ (§4).

Limitations

What we do not claim. We do not claim success equivalence between BF16 and W4 agents; equivalence tests at ± 3 pp on new-task intervals fail. G, H and F are oracle quantities, not the value of a deployable policy (§6). We do not claim that BF16 and a second W4 checkpoint are equivalent replacements. We do not claim a W$\times$KV interaction, and we do not present static precision restoration as a method (Appendices G and G.1). We do not claim a distinct action–continuation interaction mechanism, or that outcomes are mainly determined by the action: the outcome classes and the label-permutation reference (Appendix B.2) are exploratory and post hoc, with headline shares from the retail and airline development panel, and the action-over-continuation gap comes only from the 13.8% of W4 divergent states with exactly two successful arms.

Two model families, limited coverage. Most crossed forks use Qwen3 models: 4B and 8B throughout, plus a single Qwen3-14B configuration on BFCL. The second family, Mistral-Small-3.2-24B (GPTQ-W4, RTN-W4), was forked only at ordinary BF16-sourced τ^2-bench states of an 83-task conditional population (§7); it has no first tool-structure-divergence forks, no same-precision reference and no outcome-class analysis. Effects may differ at other scales and in other families. End-to-end runs also exist for Qwen3-14B and Qwen3-32B; apart from that one Qwen3-14B configuration they have no crossed forks, and we draw no conclusion from them (Appendix A.2). A time-boxed Llama-3.1-8B-Instruct check (GPTQ-W4, BFCL multi-turn) was stopped at its gate: BF16 solved 5 of 120 gate items, so the four arms carried almost no outcome variance, and it yields no evidence for that family.

Simulated quantization. Quantized checkpoints are dequantized and served with BF16 compute on a deterministic path. Deployed low-bit kernels change the numerics and add serving variation that we deliberately remove. We make no claims about latency, memory or deployed kernels, and the cost of a BF16 fallback is outside our scope.

One user simulator; two small controls. All τ^2-bench episodes use a single greedy simulator, and continuation effects include its replies. Both second-simulator controls replace only the user after a fixed history, cover 78 and 80 states, and say nothing about real users. The Bank P control has 20 states per cell, too few to confirm or refute few-pp contrasts. Neither control can be reproduced draw for draw. The replacement user is a hosted Sonnet 5.5 model sampled with its default settings, without temperature or seed control. In the Bank U control, Claude Code subagents answered many relay requests in one context, so a responder could see other arms or replicates of the same conversation. The Bank P control removed this exposure by issuing one isolated, single-request call per user turn. The same-precision reference (§5) covers q4 and q8 GPTQ-W4 on the development panel (six blocks: the frozen rule takes the longest complete common prefix, and block f07 of the same-precision q8 design is incomplete) and GPTQ-W3 on three blocks, and swaps one calibration sample for another; it is our only reference for how much turnover and continuation dependence a change of model without a change of bit width produces, and it says nothing about AWQ, KV8, RTN or rotated KV4.

Small confirmation and partial banks. The pre-registered holdout covers 34 telecom tasks and three cells. Its pre-registered readings had no negative category, so a clearly negative estimate could only be read as inconclusive; its negative point estimate, inconclusive by those readings, applies only to telecom, these three cells and this simulator (its replay acceptance was audited post hoc and is complete; Appendix F). Most development results come from retail and airline, so the holdout confounds a domain shift (telecom) with generalisation to new tasks. On 80 telecom development tasks, the two q8 holdout cells show no detectable net effect (Appendix E.5), so the telecom domain alone does not produce a detectable drop; but their 6-block intervals are wide and q4 has no telecom development forks, so the holdout drop may reflect heterogeneity among its 34 tasks as much as a domain effect. Bank P intervenes only at the first tool-structure divergence per task and block. Bank U samples one state per task and block; its main analyses use three blocks, so their block intervals rest on two degrees of freedom, and the later panels use two to six blocks. Prompt nonces define replicate blocks; they are a researcher-chosen perturbation, not a deployment distribution.

Ethics Statement

This work involves no crowdworkers, no human annotation and no personal data. Where a user is present (τ^2-bench), it is simulated by an LLM; tasks are public benchmark tasks, and all agents are openly released Qwen3, Mistral-Small-3.2 and (gate check only) Llama-3.1 checkpoints, or simulated-quantized versions of them. Outcomes are the benchmarks' automatic success checks, and action categories are assigned by fixed rules; reported values map to committed analysis outputs (Appendix G.1). The main risk we see is over-reading. A small forked-state gain on one simulator and a fixed task pool does not certify a quantized agent for deployment, and our end-to-end comparisons do not establish equivalence. Deployment decisions in consequential settings should not rest on these measurements alone.

References

Victor Barres, Honghua Dong, Soham Ray, Xujie Si, and Karthik Narasimhan. 2025. τ^2-Bench: Evaluating conversational agents in a dual-control environment. *arXiv preprint arXiv:2506.07982*.

Bjarni Haukur Bjarnason, André Silva, and Martin Monperrus. 2026. On randomness in agentic evals. *arXiv preprint arXiv:2602.07150*.

Yihang Chen, Pin Qian, Su Wang, Chong Peng, Huan Xu, Xiyang Wu, and Yiqi Sun. 2026. Noise floor audit for agent benchmarks. *arXiv preprint arXiv:2608.22331*.

Peijie Dong, Zhenheng Tang, Xiang Liu, Lujun Li, Xiaowen Chu, and Bo Li. 2025. Can compressed LLMs truly act? An empirical evaluation of agentic capabilities in LLM compression. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*.

Gaoyuan Du, Anam Nawaz Khan, Rex Zhou, Xiaoyang Liu, Deepayan Chakrabarti, Fnu Suya, and Xueping Li. 2026. Greedy decoding is not Precision-Invariant: Cross-Precision output divergence in LLM inference. *Transactions on Machine Learning Research*.

Abhinav Dutta, Sanjeev Krishnan, Nipun Kwatra, and Ramachandran Ramjee. 2024. Accuracy is not all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Elias Frantar, Saleh Ashkboos, Torsten Hoefler, and Dan Alistarh. 2023. GPTQ: Accurate post-training quantization for generative pre-trained transformers. In *International Conference on Learning Representations (ICLR)*.

Ashritha Gonuguntla. 2026. The replay gap: Static evaluation of model switching in LLM agents scores the wrong world. *arXiv preprint arXiv:2608.08239*.

Horace He and Thinking Machines Lab. 2025. Defeating nondeterminism in LLM inference. Thinking Machines Lab: Connectionism (blog). September 10, 2025.

Ashish Jain and Armaan Sandhu. 2026. UserProxyBench: Evaluating LLM user simulators for agent benchmarks and training. In *Agentic AI Benchmarks and Applications for Enterprise Tasks Workshop (AABA4ET) at NeurIPS 2026*.

Jiwon Jang, Kisu Yang, Heuiseok Lim, and Hyunwoo Park. 2026. Flat score, amplified failures: How the error budget masks damage in quantized LLM agents. *arXiv preprint arXiv:2607.27275*.

Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. 2023. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*.

Mukai Li, Qingcheng Zeng, Tianqing Fang, Zhenwen Liang, Linfeng Song, Qi Liu, Haitao Mi, and Dong Yu. 2026a. Verified critical step optimization for LLM agents. In *Findings of the Association for Computational Linguistics: ACL 2026*.

Yuanhao Li, Hongbo Wang, Xuhong Chen, Yiming Cao, and Xunzhu Tang. 2026b. Counterfactual rollout replay: Forkable environments as free process rewards for software engineering agents. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Yuanzhe Li, Jianing Deng, Jingtong Hu, Tianlong Chen, Song Wang, and Huanrui Yang. 2026c. Dynamic mixed-precision routing for efficient multi-step LLM interaction. *Preprint*, arXiv:2602.02711.

Ji Lin, Jiaming Tang, Haotian Tang, Shang Yang, Wei-Ming Chen, Wei-Chen Wang, Guangxuan Xiao, Xingyu Dang, Chuang Gan, and Song Han. 2024. AWQ: Activation-aware weight quantization for LLM compression and acceleration. In *Proceedings of Machine Learning and Systems (MLSys)*.

Zirui Liu, Jiayi Yuan, Hongye Jin, Shaochen Zhong, Zhaozhuo Xu, Vladimir Braverman, Beidi Chen, and Xia Hu. 2024. KIVI: A tuning-free asymmetric 2bit quantization for KV cache. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*.

Sanae Lotfi, Polina Kirichenko, Steven Li, and Zechun Liu. 2026. Quantized reasoning models think they need to think longer, but they do not. *arXiv preprint arXiv:2606.00206*.

Haiquan Lu, Zigeng Chen, Gongfan Fang, Xinyin Ma, and Xinchao Wang. 2026. Mix-Quant: quantized prefilling, precise decoding for agentic LLMs. *arXiv preprint arXiv:2605.20315*.

Mistral AI. 2025. Mistral-Small-3.2-24B-Instruct-2506. `https://huggingface.co/mistralai/Mistral-Small-3.2-24B-Instruct-2506`. Model card.

Shishir G. Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. 2025. The Berkeley function calling leaderboard (BFCL): From tool use to agentic evaluation of large language models. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, volume 267 of *Proceedings of Machine Learning Research*.

Aditi Patodiya. 2026. Same request, different answer: Quantization amplifies Cache-Induced divergence in LLM serving. *arXiv preprint arXiv:2609.04748*.

Qwen Team. 2025. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*.

Preethi Seshadri, Samuel Cahyawijaya, Ayomide Odumakinde, Sameer Singh, and Seraphina Goldfarb-Tarrant. 2026. Lost in simulation: LLM-simulated users are unreliable proxies for human users in agentic evaluations. *arXiv preprint arXiv:2601.17087*.

Jaineet Shah. 2026. Causal agent replay: Counterfactual attribution for LLM-agent failures. *arXiv preprint arXiv:2606.08275*.

Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew Hausknecht. 2021. ALFWorld: Aligning text and embodied environments for interactive learning. In *International Conference on Learning Representations (ICLR)*.

Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. 2024. τ-bench: A benchmark for tool-agent-user interaction in real-world domains. *arXiv preprint arXiv:2406.12045*.

Jiayi Yuan, Hao Li, Xinheng Ding, Wenya Xie, Yu-Jhe Li, Wentian Zhao, Kun Wan, Jing Shi, Xia Hu, and Zirui Liu. 2025. Understanding and mitigating numerical sources of nondeterminism in LLM inference. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Guibin Zhang, Junhao Wang, Junjie Chen, Wangchunshu Zhou, Kun Wang, and Shuicheng Yan. 2025. AgenTracer: Who is inducing failure in the LLM agentic systems? *arXiv preprint arXiv:2509.03312*.

Manyi Zhang, Ji-Fu Li, Zhongao Sun, Xiaohao Liu, Zhenhua Dong, Xianzhi Yu, Haoli Bai, and Xiaobo Xia. 2026. QuantClaw: precision where it matters for OpenClaw. *arXiv preprint arXiv:2604.22577*.

A Evaluation Contract and Scope

Splits. The development split has 58 retail, 27 airline and 80 telecom tasks. The holdout has 34 telecom tasks, and its pre-registration hash is recorded in the split file. The telecom split was fixed from τ^2-bench's 114-task base split by a stratified, salted hash before any run. The external benchmarks use BFCL multi-turn (base, miss_func, miss_param and long_context, 197 items each) and ALFWorld (134 out-of-distribution games). Items were selected by hash in manifests committed before any model run.

Success endpoints. For retail and airline we use the final-database match with normal termination. This is not τ^2-bench's official reward, which additionally scores natural-language or communicate checks. Telecom uses the official deterministic task-basis success. Context overflow and unparseable output count as failures. Infrastructure outages are missing, never failures, and a block with a missing unit leaves the primary analysis.

Fork mechanics. The forced action is returned by a fork agent for the one step that follows h_t. A tool call is executed by the real environment, and text is answered by the real simulator, which never reuses observations from the original trajectory. τ^2-bench resets step and error counters when it starts from a history. The fork restores them to the values used at t, so the remaining budget equals the original. Every max_steps and too_many_errors episode on the checked blocks is reproduced by the recount.

Replay acceptance details. In the engineering gate, one B→B field, the official reward, was missing from older trajectories and was recomputed from the stored reward information before the verdict. Besides the 924 sampled development states, 328/328 sampled Bank U states passed the QQ no-op, and the Bank U gate passed a further 135/135 no-ops (a different set from the holdout's 135 QQ no-ops, all of which also passed; Appendix F). Serving in separate processes closes the prefix-cache channel (Patodiya, 2026). The same-precision run's acceptance counts are in Appendix C.

Determinism of the serving stack. Closed-loop determinism pairs, the same configuration on two nodes at different times, reproduced whole episodes token for token for dense BF16 paths at

tensor parallelism 1 and 2 with 1 to 16 concurrent requests (e.g. 257/257 and 1520/1520 episodes). Nine arms that two studies ran independently at the same states were byte-identical. We did not test exactness under KV-cache preemption, or on other models, GPUs or vLLM versions. A deployed AWQ kernel path with fp16 dequantization was not replay-exact, and no reported outcome uses it.

User simulator. The served user simulator is Qwen3-30B-A3B-Instruct in every episode; a user-model label stored in the fork state files is outdated and should be ignored.

A.1 Separating local value from end-to-end success

The replacement gains G_{all}, G, H and F are not end-to-end comparisons. Table 2 gives the end-to-end comparison, computed separately on the same checkpoints, the same tasks and the same 20 nonce blocks as the forks.

Cell	Q / BF16 (%)	Q−BF16 [95% CI]	NI 3 pp	TOST ±3 pp
q4 GPTQ-W4	40.0 / 35.8	+4.2 [−2.4, +10.5]	pass	fail
q4 AWQ-W4	36.8 / 35.8	+0.9 [−5.8, +7.5]	fail	fail
q8 GPTQ-W4	30.1 / 30.4	−0.2 [−7.7, +6.8]	fail	fail
q8 AWQ-W4	31.5 / 30.4	+1.2 [−5.3, +7.6]	fail	fail
pool, all four		+1.5 [−2.7, +5.6]	pass	fail

Table 2: End-to-end closed-loop success, quantized minus BF16 (pp). Units are paired on nonce, and the crossed task×block bootstrap is drawn jointly across cells. Non-inferiority (NI) at a 3 pp margin passes for q4 GPTQ-W4 and the all-four pool. Equivalence (TOST at ±3 pp) passes nowhere. These comparisons therefore support neither a difference nor equivalence.

A.2 Coverage Matrix

Table 3 lists the completed crossed-intervention coverage behind every Qwen3 Bank P and Bank U result; the coverage of the U2B panels (§6, §7) is in Appendix D. It counts tasks, states and complete blocks after the strict admission rule. Four of the five stress cells turn over 18.4–20.4% of units with three blocks; we keep them separate from W4, and no W2 crossed-intervention panel is reported.

B Which Upgrade Changes the Task?

Table 4 gives every cell of Figure 1b, including the stress cells, with new-task intervals. Table 5 gives every cell of Figure 4a, with $F - G$, $F - H$ and the $(\Delta_{\text{Q}}, \Delta_{\text{B}})$ summaries.

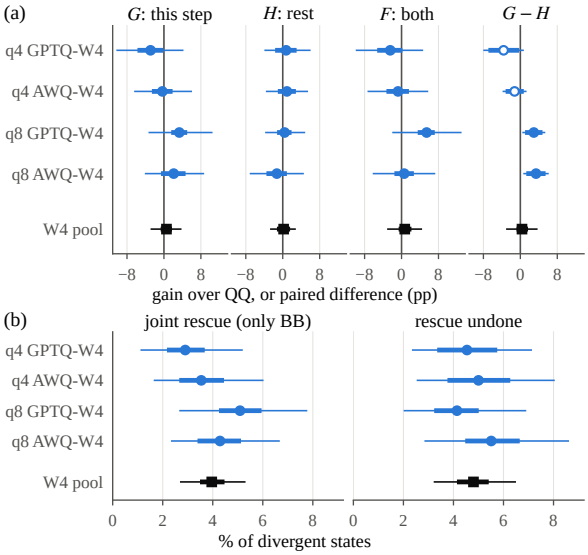


Figure 4: Upgrade modes on the same divergent states (τ^2-bench development; four core W4 cells and their pool). (a) Replace this step (G), take over from the next step (H), or both (F), as gain in success over QQ, and the paired difference $G - H$; pp; unadjusted, thick: fixed-pool 95%, thin: new-task 95%, except for the cells' $G - H$ (the Holm family), where thin is the simultaneous interval and only filled markers are significant after Holm adjustment. (b) Joint rescue (only BB succeeds) and rescue undone ($Y_{\text{BQ}} = 1$, $Y_{\text{QQ}} = 0$, $Y_{\text{BB}} = 0$), % of divergent states (exploratory, post hoc); thick: fixed-pool 95%, thin: new-task 95%. Each is a rescue that one single-mode evaluation misses. Interval details: below.

B.1 Paired comparisons and interval families

Paired comparisons distinguish upgrade modes. All three upgrade modes are measured on the same states, so they can be compared by paired differences (Figure 4a). Over the four core W4 cells, $G - H$ is a fixed Holm family ($m = 4$; fixed task pool, t over blocks). In both Qwen3-8B cells, action replacement outperforms continuation takeover: $G - H = +2.9$ (GPTQ-W4) and $+3.4$ pp (AWQ-W4), Holm-adjusted $p = .013$ and $.012$; this ranks the two modes and does not show that replacing the action beats QQ. In the Qwen3-4B cells the point estimates favour continuation takeover (-3.6 and -1.3 pp), but neither is significant after Holm adjustment. All four new-task intervals include zero, so the finding is limited to the fixed task pool and stays exploratory (intervals and further cautions below). In the low-success stress cell q8 RTN-W3, continuation takeover beats action replacement ($G - H = -15.7$ [-27.5, -3.9] pp, three blocks). For Qwen3-14B GPTQ-W4 on BFCL (one ex-

Cell	τ^2-bench retail (Bank P)	τ^2-bench airline (Bank P)	telecom dev (Bank P)	Bank U (retail / airline)	BFCL / ALFWorld
q4 GPTQ-W4	58 t · 727 st · 20/20 b	27 t · 375 st · 20/20 b	holdout only	3 b · 393 / 177 rows	788 items · 573 st / –
q4 AWQ-W4	58 t · 986 st · 20/20 b	27 t · 453 st · 20/20 b	–	3 b · 393 / 177 rows	788 items · 642 st / –
q4 KV8	58 t · 412 st · 22/22 b	27 t · 181 st · 22/22 b	–	3 b · 393 / 177 rows	– / –
q8 GPTQ-W4	58 t · 1017 st · 20/20 b	27 t · 456 st · 20/20 b	80 t · 477 st · 6/6 b; holdout	3 b · 393 / 177 rows	788 · 533 st / 134 · 115 st
q8 AWQ-W4	58 t · 1087 st · 20/20 b	27 t · 475 st · 20/20 b	80 t · 480 st · 6/6 b; holdout	3 b · 393 / 177 rows	788 · 548 st / 134 · 119 st
q8 KV8	58 t · 585 st · 22/22 b	27 t · 270 st · 22/22 b	–	3 b · 393 / 177 rows	– / –
q4 GPTQ-W3	58 t · 162 st · 3/3 b	27 t · 70 st · 3/3 b	–	–	591 · 540 st† / –
q4 KV4-rot	58 t · 164 st · 3/3 b	27 t · 69 st · 3/3 b	–	–	– / –
q8 GPTQ-W3	58 t · 166 st · 3/3 b	27 t · 76 st · 3/3 b	–	–	788 · 696 st / 134 · 130 st
q8 KV4-rot	58 t · 165 st · 3/3 b	27 t · 74 st · 3/3 b	–	–	– / –
q8 RTN-W3	58 t · 174 st · 3/3 b	27 t · 81 st · 3/3 b	–	–	– / –
q4 / q8 KV4 (unrotated)	–	–	–	–	788 · 787 / 788 · 788 st; q8 ALFWorld 134 · 134 st
q14 GPTQ-W4	–	–	–	–	788 · 471 st / –

Table 3: Completed crossed-intervention coverage (t = tasks, st = states, b = complete/found blocks; BFCL and ALFWorld: items · states, one run). Bank U: blocks completed and state rows per domain; blocks f01–f03 are analysed (Appendix E.2). "holdout": the cell is in the telecom holdout (Appendix F). Qwen3-32B has end-to-end runs only; Llama-3.1-8B stopped at its fork gate (Limitations); the q14 (Qwen3-14B) BFCL row is a within-family scale check. q4 RTN-W3 has no complete Bank P block and is omitted. The telecom development results (Appendix E.5) use the 6 blocks listed. †The BFCL base category of q4 GPTQ-W3 was dropped by the QQ no-op rule (Appendix E.4). The end-to-end q4 AWQ-W3 run is incomplete (two airline episodes hit 600 s serving timeouts), and so are the q4 RTN-W2 stress runs (context-window overflows of the degenerate model and a truncated result file); neither is reported.

Cell	units	div.	r	d	M	G_{all} [fixed pool]	[new tasks]
q4 GPTQ-W4	85×20	1102	6.8	8.7	15.5	-1.9 [-3.7, -0.1]	[-6.8, $+2.7$]
q4 AWQ-W4	85×20	1439	9.8	10.1	19.9	-0.3 [-2.2, $+1.6$]	[-5.4, $+5.2$]
q8 GPTQ-W4	85×20	1473	10.4	7.5	17.9	$+2.9$ [$+1.4$, $+4.4$]	[-2.9, $+9.2$]
q8 AWQ-W4	85×20	1562	11.4	9.5	20.9	$+1.9$ [-0.6, $+4.3$]	[-3.9, $+8.1$]
W4 pool	85×20		9.6	9.0	18.6	$+0.7$ [-0.3, $+1.6$]	[-2.1, $+3.5$]
q4 KV8	85×22	593	3.1	3.4	6.5	-0.3 [-1.7, $+1.0$]	[-3.0, $+2.3$]
q8 KV8	85×22	855	4.3	3.7	8.1	$+0.6$ [-0.4, $+1.6$]	[-1.9, $+3.0$]
q4 GPTQ-W3^s	85×3	232	9.0	10.2	19.2	-1.2 [-7.0, $+4.7$]	[-8.6, $+5.9$]
q4 KV4-rot^s	85×3	233	11.8	7.1	18.8	$+4.7$ [-4.1, $+13.5$]	[-3.1, $+12.9$]
q8 GPTQ-W3^s	85×3	242	10.2	8.2	18.4	$+2.0$ [-12.5, $+16.4$]	[-6.3, $+11.8$]
q8 KV4-rot^s	85×3	239	11.8	8.6	20.4	$+3.1$ [-14.5, $+20.8$]	[-5.9, $+12.9$]
q8 RTN-W3^s	85×3	255	3.1	3.5	6.7	-0.4 [-3.8, $+3.0$]	[-4.7, $+3.9$]

Table 4: Replacing Q's first tool-structure-divergent action while Q continues, τ^2-bench retail/airline panel (pp). Denominator: the full pool of tasks × complete blocks; units without a Bank P state count as zero. div. = Bank P states. $G_{\text{all}} = r - d$ is a local replacement effect, *not* the end-to-end success difference (Appendix A.1). Fixed pool: block bootstrap (≥10 blocks) or t over block means; new tasks: crossed task×block bootstrap. ^{s}stress cell, 3 blocks.

ploratory configuration), action replacement wins ($G - H = +5.3$ [$+2.3$, $+8.4$] pp), and action and continuation are complementary there (below). The joint upgrade F has the largest point estimate in both examples, and we compare no compute costs, so no cost-optimal mode is established.

The $G - H$ family and Figure 4a. In Figure 4a, G, H, F and the pool show fixed-pool 95% (thick; block bootstrap) and new-task 95% (thin; crossed) intervals, none multiplicity-adjusted. For the four cells' $G - H$, which form the Holm family ($m = 4$), thick is the plain 95% t interval and thin the Bonferroni simultaneous 98.75% interval; filled markers are significant after Holm adjustment, open markers are not. Above, the stress cell's $G - H$ has a t interval over three blocks, and the BFCL configuration's an item-bootstrap 95% interval; neither is drawn. The plain 95% t intervals are [$+1.0$, $+4.8$] (q8 GPTQ-W4, $G - H = +2.9$, Holm-adjusted $p = .013$) and [$+1.3$, $+5.5$] (q8 AWQ-W4, $+3.4$, $p = .012$). In q4 GPTQ-W4 (-3.6) the plain interval [-6.9, -0.3] excludes zero, but the Holm-adjusted $p = .069$; q4 AWQ-W4 (-1.3) has $p = .199$. Two further cautions apply. A significant result in one model and a non-significant one in the other is not a significant difference between models, and the cells select different divergent states. $G - H > 0$ shows that replacing the action beats handing over the continuation, not that replacing the action beats QQ.

Supplementary configuration. For Qwen3-14B GPTQ-W4 on BFCL (one configuration, one run per item, 471 divergent items; exploratory), action and continuation are complementary: keeping Q's action, a BF16 continuation adds $H = +0.4$ [-1.1, $+1.9$] pp; after B's action has replaced Q's, the BF16 continuation adds a further $F - G = +3.4$ [$+1.3$, $+5.6$] pp.

Table 6 gives all new-task intervals corresponding to Table 5.

B.2 Outcome Classes and a Label-Permutation Reference

Everything in this section is exploratory and post hoc. It uses the four outcomes of each divergent state in the order BB, BQ, QB, QQ; k is the num-

Cell	states	G (this step)	H (next step on)	F (from here on)	$G-H$	$F-G$	$F-H\ (=\Delta_{\mathrm{B}})$	off-diag. %	rev. %	mean I
q4 GPTQ-W4	1102	-2.9 [-5.7, -0.2]	$+0.7$ [-1.6, $+3.0$]	-2.5 [-5.3, $+0.2$]	-3.6 [-6.6, -0.6]	$+0.5$ [-1.2, $+2.2$]	-3.2 [-6.4, $+0.2$]	25.0	2.6	$+0.3$
q4 AWQ-W4	1439	-0.3 [-2.6, $+1.9$]	$+0.9$ [-1.1, $+2.9$]	-0.8 [-3.3, $+1.6$]	-1.3 [-3.0, $+0.6$]	-0.4 [-2.4, $+1.5$]	-1.7 [-3.6, $+0.4$]	29.3	3.1	$+1.3$
q8 GPTQ-W4	1473	$+3.3$ [$+1.6$, $+5.0$]	$+0.4$ [-1.2, $+1.9$]	$+5.4$ [$+3.5$, $+7.2$]	$+2.9$ [$+1.3$, $+4.7$]	$+2.1$ [$+0.2$, $+4.0$]	$+5.0$ [$+3.2$, $+6.8$]	22.9	2.7	-1.7
q8 AWQ-W4	1562	$+2.1$ [-0.6, $+4.7$]	-1.3 [-3.5, $+0.9$]	$+0.6$ [-1.5, $+2.7$]	$+3.4$ [$+1.5$, $+5.2$]	-1.5 [-4.2, $+0.8$]	$+1.9$ [-0.1, $+3.8$]	30.5	4.2	$+0.3$
W4 pool	5576	$+0.5$ [-0.7, $+1.7$]	$+0.2$ [-1.3, $+1.6$]	$+0.7$ [-0.6, $+2.1$]	$+0.4$ [-0.8, $+1.7$]	$+0.2$ [-1.0, $+1.2$]	$+0.5$ [-1.0, $+2.1$]	26.9	3.2	$+0.0$
q4 KV8	593	-1.0 [-5.3, $+3.2$]	-2.5 [-4.5, -0.5]	-1.2 [-4.7, $+2.3$]	$+1.5$ [-3.1, $+6.0$]	-0.2 [-3.5, $+3.7$]	$+1.3$ [-2.5, $+4.9$]	15.7	1.2	-2.4
q8 KV8	855	$+1.3$ [-0.9, $+3.4$]	-0.4 [-2.1, $+1.5$]	$+2.5$ [$+0.0$, $+5.1$]	$+1.6$ [-0.6, $+3.9$]	$+1.2$ [-0.5, $+3.2$]	$+2.8$ [$+0.3$, $+5.5$]	14.9	1.2	-1.5
KV8 pool	1448	$+0.1$ [-2.8, $+3.0$]	-1.4 [-2.7, -0.1]	$+0.6$ [-1.8, $+3.1$]	$+1.6$ [-1.2, $+4.1$]	$+0.5$ [-1.6, $+2.8$]	$+2.1$ [-0.2, $+4.2$]	15.3	1.2	-1.9
q4 GPTQ-W3^{s}	232	-1.3 [-7.7, $+5.0$]	$+4.3$ [-10.7, $+19.4$]	$+1.3$ [-1.9, $+4.4$]	-5.6 [-15.0, $+3.7$]	$+2.6$ [-0.6, $+5.8$]	-3.0 [-15.1, $+9.1$]	31.5	3.9	$+1.7$
q4 KV4-rot^{s}	233	$+5.2$ [-4.5, $+14.9$]	$+9.4$ [$+5.1$, $+13.7$]	$+12.4$ [$+11.0$, $+13.9$]	-4.3 [-18.3, $+9.8$]	$+7.3$ [-2.0, $+16.5$]	$+3.0$ [-1.9, $+8.0$]	33.0	1.7	$+2.1$
q8 GPTQ-W3^{s}	242	$+2.1$ [-13.3, $+17.2$]	$+0.4$ [-7.4, $+8.2$]	$+2.1$ [-7.4, $+11.4$]	$+1.7$ [-9.1, $+12.3$]	$+0.0$ [-11.0, $+11.1$]	$+1.7$ [-0.1, $+3.4$]	23.1	4.5	$+0.4$
q8 KV4-rot^{s}	239	$+3.3$ [-15.4, $+21.9$]	-0.8 [-18.2, $+16.3$]	$+4.6$ [-16.0, $+25.0$]	$+4.2$ [$+2.5$, $+5.8$]	$+1.3$ [-1.8, $+4.3$]	$+5.4$ [$+0.8$, $+10.0$]	31.0	4.6	-2.1
q8 RTN-W3^{s}	255	-0.4 [-3.8, $+3.0$]	$+15.3$ [$+0.7$, $+29.9$]	$+18.4$ [$+16.7$, $+20.1$]	-15.7 [-27.5, -3.9]	$+18.8$ [$+13.8$, $+23.9$]	$+3.1$ [-13.0, $+19.2$]	29.0	0.4	-3.5

Table 5: Three replacement gains and their paired differences on each cell's Bank P states (pp; same denominator in every column). $G = Y_{\mathrm{BQ}} - Y_{\mathrm{QQ}}$, $H = Y_{\mathrm{QB}} - Y_{\mathrm{QQ}}$, $F = Y_{\mathrm{BB}} - Y_{\mathrm{QQ}}$. Off-diag.: share of states with $\Delta_{\mathrm{Q}} \neq \Delta_{\mathrm{B}}$; rev.: strict reversals. Fixed-pool 95% intervals (block bootstrap over 20–22 complete blocks; t over 3 blocks for ^{s}stress cells, which cannot be pooled). New-task intervals are in Table 6. Within a row, contrasts are paired and identified; between rows, descriptive.

Cell	G	H	F	$G-H$	$F-G$	$F-H$
q4 GPTQ-W4	[-10.3, $+4.2$]	[-4.0, $+6.0$]	[-9.9, $+4.7$]	[-11.2, $+3.6$]	[-4.7, $+6.1$]	[-11.1, $+4.4$]
q4 AWQ-W4	[-6.5, $+6.1$]	[-3.7, $+5.5$]	[-7.4, $+5.8$]	[-7.7, $+5.1$]	[-6.2, $+5.4$]	[-8.7, $+4.7$]
q8 GPTQ-W4	[-3.3, $+10.5$]	[-3.8, $+4.9$]	[-2.0, $+13.0$]	[-3.2, $+9.3$]	[-3.0, $+7.1$]	[-2.0, $+12.5$]
q8 AWQ-W4	[-4.1, $+8.7$]	[-7.1, $+4.6$]	[-6.2, $+7.3$]	[-1.6, $+8.6$]	[-7.6, $+4.2$]	[-3.8, $+7.5$]
W4 pool	[-2.9, $+3.8$]	[-2.7, $+2.8$]	[-3.1, $+4.4$]	[-3.1, $+3.7$]	[-2.8, $+3.1$]	[-3.5, $+4.5$]
q4 KV8	[-9.4, $+7.3$]	[-6.4, $+1.1$]	[-9.3, $+7.7$]	[-7.1, $+9.7$]	[-5.7, $+5.3$]	[-6.7, $+9.5$]
q8 KV8	[-4.0, $+6.9$]	[-3.5, $+2.9$]	[-3.5, $+8.1$]	[-3.8, $+7.2$]	[-2.5, $+5.0$]	[-3.1, $+8.7$]
KV8 pool	[-5.3, $+5.6$]	[-3.7, $+0.9$]	[-4.5, $+5.8$]	[-3.8, $+6.6$]	[-2.8, $+3.9$]	[-3.0, $+7.3$]
q4 GPTQ-W3^{s}	[-9.7, $+6.6$]	[-6.6, $+14.5$]	[-6.8, $+9.4$]	[-15.0, $+4.0$]	[-5.2, $+10.5$]	[-12.2, $+6.6$]
q4 KV4-rot^{s}	[-3.4, $+14.2$]	[$+2.1$, $+17.0$]	[$+3.0$, $+21.6$]	[-13.5, $+7.0$]	[-1.3, $+17.1$]	[-6.2, $+12.2$]
q8 GPTQ-W3^{s}	[-6.7, $+12.1$]	[-7.2, $+8.1$]	[-6.6, $+11.8$]	[-7.7, $+11.2$]	[-9.1, $+9.1$]	[-6.8, $+11.0$]
q8 KV4-rot^{s}	[-6.6, $+13.8$]	[-11.2, $+8.5$]	[-6.6, $+15.9$]	[-3.3, $+12.2$]	[-7.2, $+9.9$]	[-2.5, $+14.6$]
q8 RTN-W3^{s}	[-4.7, $+3.9$]	[$+5.9$, $+25.9$]	[$+9.4$, $+28.2$]	[-24.7, -7.1]	[$+10.2$, $+27.8$]	[-7.1, $+13.3$]

Table 6: New-task generalisation: crossed task×block bootstrap 95% intervals for the estimates of Table 5 (pp; same Bank P states). These answer whether the effect would hold for new tasks drawn like the panel; the fixed-pool intervals of Table 5 treat the 85 tasks as fixed. ^{s}stress cell, 3 blocks.

ber of successful arms.

Five classes. The 16 patterns fall into five classes: unchanged (`0000`, `1111`); action only (`1100`, `0011`: the outcome follows the action under both continuations); continuation only (`1010`, `0101`); one arm differs from the other three (k odd); strict reversal (`1001`, `0110`). Changing the action changes the outcome under at least one continuation exactly in the action-only, one-arm and reversal classes, and changing the continuation exactly in the continuation-only, one-arm and reversal classes. Both change it exactly on the off-diagonal states, so the two shares overlap and must not be added, and their gap equals action only minus continuation only: it comes from $k = 2$ states alone and restates the $k = 2$ arrangement below on all divergent states. Two named patterns follow the upgrade modes: joint rescue (`1000`: only the joint upgrade succeeds) and rescue undone ($Y_{\mathrm{BQ}} = 1$, $Y_{\mathrm{QQ}} = 0$, $Y_{\mathrm{BB}} = 0$: B's action rescues under Q's continuation and B's takeover of the continuation undoes it). Table 7 gives the shares. The action-side gap holds in the τ^2-bench retail/airline W4 and KV8 pools, in the W4 same-precision fork and in the BFCL W4 pool. It is near zero on the telecom development tasks ($+0.1$ [-2.1, $+2.4$] pp, 6-block t) and in the q4 GPTQ-W3 and KV4-rot stress cells, and it is not resolved in the ALFWorld W4 pool ($+3.4$ [-1.8, $+8.4$]) or the W3 same-precision pair ($+1.9$ [-2.9, $+6.7$]). Its point estimate is negative in the q8 RTN-W3 stress cell (-7.8; 3-block t [-18.1, $+2.4$], new-task [-15.3, -1.2]) and, not shown, in ALFWorld q8 GPTQ-W3 (-19.2 [-27.3, -10.8]). In the ALFWorld and BFCL KV4 cells (not shown) Q almost never succeeds, so continuation-only states dominate there mechanically.

W4 pool details. In the τ^2-bench W4 pool, the gap between the action and continuation shares is $+5.2$ pp (new-task [$+3.1$, $+7.7$]): action effects are somewhat more common, and the continuation changes the outcome mostly where the action matters too. Of the 26.9% off-diagonal states, 23.8% have one arm differing from the other three and 3.2% are strict reversals; states where the swap helps more under Q ($I > 0$, 13.6%) and under B

($I < 0$, 13.4%) are about equally frequent, so the mean interaction is $+0.0$ [-1.9, $+2.0$] pp. With fixed-pool, then new-task 95% intervals, changing the action changes the outcome under at least one continuation in 34.9% ([33.6, 36.0]; [28.8, 41.0]) of W4 divergent states, and changing the continuation controller does so under at least one action in 29.6% ([28.5, 30.7]; [24.2, 35.3]). The two sets overlap in the 26.9% ([25.7, 28.1]; [22.1, 32.0]) where the swap's effect depends on the continuation, so the shares must not be added; all four arms agree in 62.4%.

A label-permutation reference. Given k, assigning the four outcomes to the four arms uniformly at random gives the off-diagonal share $E_0 = P(k \text{ odd}) + P(k = 2)/3$. In every cell the observed share minus E_0 equals $P(k = 2)\,[P(\text{reversal} \mid k = 2) - 1/3]$, so it can depart from E_0 only through strict reversals. In the W4 pool the share is 26.9% against $E_0 = 28.4\%$ (-1.4 [-1.8, -1.0] pp, fixed pool, block bootstrap over 20 blocks; KV8: 15.3 against 18.9%), because reversals are 23.6% [19.2, 28.3] (task bootstrap) of $k = 2$ states instead of one third. This shows only that strict reversals are rarer than under random labelling. It does not show that the continuation controller is irrelevant, which would require $Y_{\mathrm{BB}} = Y_{\mathrm{BQ}}$ and $Y_{\mathrm{QB}} = Y_{\mathrm{QQ}}$, or that the interaction is zero: the reference itself allows $I \neq 0$ in most arrangements. Conversely, an off-diagonal state does not show a distinct interaction mechanism: if an arm succeeds when an action score plus a continuation score exceeds a threshold, exactly additive scores can make only BB succeed (`1000`), although they never produce a strict reversal.

Joint rescue and rescue undone against the reference. The same reference, applied to the two named patterns, keeps k for every state and permutes the four outcome labels uniformly: `1000` then has probability $1/4$ at $k = 1$, and rescue undone (`0100` or `0110`) has $1/4$ at $k = 1$ plus $1/6$ at $k = 2$, so the reference shares are $P(k{=}1)/4$ and $P(k{=}1)/4 + P(k{=}2)/6$, computed on the same states and bootstrap draws as the observed shares (Table 8). In the τ^2-bench W4 pool, joint rescue exceeds the reference by $+0.47$ pp on the fixed task pool ([$+0.06$, $+0.87$]), but the new-task interval spans zero ([-0.33, $+1.37$]); rescue undone is less frequent than the reference (-0.99 [-1.58, -0.42]; new tasks [-2.19, $+0.27$]). The excess of joint rescue is concentrated in the q8 cells, and no group shows a new-task excess of rescue undone. Because the reference keeps k and does not model execution, a share near it says that the arrangement is about as common as a random arrangement of the same number of successes, not that the outcomes are noise or that the continuation is irrelevant.

BF16 and Q′ protocols. On the 1020 paired task×block units of the same-precision fork, the BF16 design detects a joint rescue on 3.92% of units and the Q′ design on 1.76%, a paired difference of $+2.16$ pp (t [$+0.34$, $+3.97$]; crossed [-0.20, $+4.71$]); for rescue undone the difference is -1.47 (t [-2.97, $+0.03$]; crossed [-3.83, $+0.78$]). Each design forks at its own first divergence, so these are differences in how often two selection-and-replacement protocols find the patterns, not precision effects at a common state; the common-state comparison (Appendix C) is the direct test of BF16's extra action value.

Mistral. On Mistral-Small-3.2-24B's own BF16-trajectory states (Appendix E.3; 464 states with differing actions, GPTQ-W4 and RTN-W4 pooled, five blocks), joint rescue is 3.4% against a reference of 2.9% and rescue undone 4.0% against 4.7%; neither difference is resolved. There BB is the recorded BF16 original rather than an upgrade of Q's state, so the denominators and the reading of the named patterns differ from Bank P, and we do not compare the two populations.

The $k = 2$ arrangement. At $k = 2$ the two successes share the action in 56.4% of W4 states and 80.3% of KV8 states, against one third under random labelling. A statement that outcomes are mainly determined by the action holds only within this subset, which is 13.8% (W4) and 14.5% (KV8) of divergent states.

C What Does Higher Precision Add?

Per model (final analysis set, $K = 6$), the turnover difference is $\Delta M = +0.6$ [-3.1, $+4.3$] pp in q4 and -2.5 [-5.6, $+0.5$] in q8 (fixed pool, 95%), with crossed intervals [-5.5, $+6.9$] and [-7.3, $+2.7$]. Within ± 3 pp, equivalence is shown for no per-model ΔM or ΔG_{all}; the q8 off-diagonal difference ($+0.40$, 90% [-2.12, $+2.92$]) is the only per-model contrast inside the margin. The protocol originally mapped "interval includes 0 and

Group	states	action changes	continuation changes	gap (pp)	unch.	act. only	cont. only	one arm	rev.	joint resc.	resc. undone
τ^2-bench W4 pool (4 cells)	5576	34.9 [33.6, 36.0]	29.6 [28.5, 30.7]	+5.2 [+4.4, +6.1]	62.4	7.9	2.7	23.8	3.2	4.0	4.8
τ^2-bench KV8 pool	1448	26.8 [25.0, 28.6]	17.0 [14.9, 19.1]	+9.8 [+8.0, +11.6]	71.5	11.6	1.7	14.1	1.2	3.0	2.3
Telecom q8 W4 pool	957	22.1 [18.2, 25.9]	21.9 [17.8, 26.1]	+0.1 [−2.1, +2.4]	74.8	3.2	3.1	18.1	0.7	3.2	4.2
Same-precision W4 (Q′ as B)	784	33.5 [29.8, 37.1]	28.7 [24.8, 32.2]	+4.9 [+2.9, +6.9]	64.0	7.4	2.5	22.5	3.6	2.3	5.8
BF16, same blocks	784	35.1 [32.6, 37.4]	27.8 [23.8, 31.6]	+7.3 [+4.9, +9.7]	62.0	10.2	2.9	22.4	2.6	5.0	3.8
q8 RTN-W3[s]	255	29.8 [28.1, 31.5]	37.6 [26.0, 49.3]	−7.8 [−18.1, +2.4]	61.6	0.8	8.6	28.6	0.4	12.6	0.8
BFCL W4 pool[b]	2296	18.4 [16.4, 20.5]	11.2 [9.8, 12.6]	+7.2 [+5.8, +8.7]	80.8	8.1	0.8	9.8	0.5	2.4	1.4
ALFWorld q8 W4 pool[b]	234	44.9 [38.3, 52.4]	41.5 [34.3, 49.4]	+3.4 [−1.8, +8.4]	50.0	8.5	5.1	33.4	3.0	9.4	5.5

Table 7: Which intervention changes the outcome, in % of divergent states (pools: equal weight per cell). "Action changes": the outcome changes with the action under at least one continuation; "continuation changes": with the continuation controller under at least one action; gap = action only − continuation only. Fixed-pool 95% intervals: block bootstrap for the 20- and 22-block τ^2-bench W4 and KV8 pools; t over per-block ratios for the 6-block telecom and same-precision rows and the 3-block [s]stress cell; [b]one run per item, item bootstrap. Same-precision rows: Q′, a second GPTQ checkpoint of the same bit width, in the B role, and BF16 on the same blocks. New-task (crossed) intervals for the τ^2-bench W4 pool: above.

		joint rescue (`1000`)				rescue undone (`0100`, `0110`)			
Group	states	obs.	ref.	obs.−ref. [fixed pool]	[new tasks]	obs.	ref.	obs.−ref. [fixed pool]	[new tasks]
τ^2-bench W4 pool	5576	3.96	3.48	+0.47 [+0.06, +0.87]	[−0.33, +1.37]	4.80	5.78	−0.99 [−1.58, −0.42]	[−2.19, +0.27]
q4 W4 cells	2541	3.22	3.56	−0.34 [−0.83, +0.13]	[−1.46, +0.88]	4.77	6.01	−1.24 [−2.09, −0.45]	[−2.97, +0.69]
q8 W4 cells	3035	4.69	3.40	+1.29 [+0.79, +1.78]	[−0.07, +2.74]	4.82	5.56	−0.74 [−1.35, −0.07]	[−2.34, +0.98]
τ^2-bench KV8 pool	1448	2.99	2.17	+0.81 [+0.10, +1.52]	[−0.51, +2.19]	2.31	4.58	−2.28 [−2.98, −1.58]	[−3.62, −0.98]
Telecom q8 W4 pool	957	3.24	3.94	−0.70 [−1.46, +0.05]	[−2.47, +1.15]	4.18	5.13	−0.95 [−1.70, −0.20]	[−2.63, +0.61]
Same-precision W4 (Q′ as B)	784	2.27	3.50	−1.23 [−2.47, −0.04]	[−2.88, +0.51]	5.77	5.75	+0.02 [−2.33, +2.25]	[−2.66, +2.92]
BF16, same blocks	784	4.96	3.53	+1.43 [−0.40, +3.42]	[−0.90, +4.26]	3.82	6.14	−2.32 [−4.46, −0.19]	[−4.68, +0.17]
Mistral, own BF16 states[m]	464	3.4	2.9	+0.5 [−2.4, +3.5]	[−2.0, +3.6]	4.0	4.7	−0.7 [−2.0, +0.7]	[−3.0, +1.9]

Table 8: Joint rescue and rescue undone against the conditional label-permutation reference (keeps the number k of successful branches per state), % of divergent states; exploratory, post hoc. 95% intervals: fixed pool (block bootstrap for the 20- and 22-block τ^2-bench pools, t over per-block ratios for fewer than 10 blocks: 6, or 5 for Mistral), new tasks (crossed task×block). Same-precision rows each fork at their own first divergence. [m]States whose actions differ, on BF16 trajectories (BB is the BF16 original), five blocks; not comparable with Bank P rows.

$|\Delta M| \leq 3$ pp" to "similar"; after the interim result had been seen, an amendment replaced that reading by two one-sided tests, which are stricter, and changed no estimate. Acceptance: every state in both designs passed Q replay and every branch passed the prefix and environment checks; QQ no-op samples passed 265/265 (precision) and 242/242 (same precision), and no directory was quarantined. At the unit level, Cohen's κ between the action-side flip indicators of the two substitutions is 0.43 in q4 and 0.61 in q8; at the task level, cross-design agreement at matched noise is 0.765 of within-design replication.

W3 stress check (§5; $K = 3$): every state passed Q replay, QQ no-op samples passed 31/31 and 39/39, and no directory was quarantined. All branches passed the prefix and environment checks except one of 774 in q8 same precision: a Q′Q′ branch that ended in an agent-side context overflow, for which the harness records no prefix comparison; its environment hash and first replayed prompt match the frozen state, and by the protocol it counts as a failure, not an exclusion.

Equivalence tests (W4). The difference in turnover is $\Delta M = M - M' = -0.98$ pp (90% CI $[-3.26, +1.30]$). No precision-specific excess was detected, but equivalence within ± 3 pp is not shown either (two one-sided tests, $p = 0.067$), and the same holds for the net gain ($\Delta G_{\text{all}} = -1.57$ $[-5.17, +2.04]$, $p = 0.230$) and, descriptively, for the off-diagonal share (-1.210 $[-4.09, +1.67]$, $p = 0.133$, as the mean of per-block ratios, which the interval and test use; -1.216 as the ratio of block sums).

Overlap details (W4). The two action-replacement flip sets share 107 units, against 76 [69, 84] under a within-task block shuffle that preserves task-level flip propensity. When B also takes over the continuation, rescue sets under the continuation-only contrast agree with $\kappa = 0.20$, and 70% of continuation-dependent units are continuation-dependent under one substitution only.

Full substitution (W4). Under full substitution B flips more units than Q′ (193 vs. 149;

+4.3 [+0.4, +8.3] pp, one of several contrasts); the action-side net difference (ΔG_{all}, B's action against Q′'s under Q's continuation, not an F difference) comes only from substitution-specific units and is not resolved (-1.6 [-6.2, $+3.0$] pp). Only value measures on the same units reveal this structure.

W3 stress pair. The protocol's complete W3 panel (Q GPTQ-W3, Q′ a second W3 calibration draw; q4 and q8, $K = 3$ blocks q3–q5, 510 units) is reported next to W4, never pooled with it. The same-precision substitution produces 17.3 pp of turnover, the BF16 substitution 18.8 pp; off-diagonal shares are 27.6% and 27.3%. $\Delta M = +1.6$ pp (95% [-0.66, $+3.80$]) is unresolved: the preregistered ± 3 pp equivalence test fails ($p = 0.055$); that the 90% interval lies just above zero does not change this. ΔG_{all} is equivalent within ± 3 pp only in the preregistered pooled population ($+1.2$, $p = 0.017$); per model it is $+3.1$ (q4) and -0.8 pp (q8). The off-diagonal difference is -0.288 (ratio of block sums) or -0.307 (mean of per-block ratios, used by the interval and test: 90% [-11.0, $+10.4$], $p = 0.27$). On the action side (G), rescued and harmed task×block units overlap ($\kappa = 0.60, 0.68$); the designs may first diverge at different decisions, so this is unit overlap, not a same-state replication. Of 122 units whose G outcome flips, 62 flip under both substitutions, 34 only under BF16 and 26 only under the other W3 checkpoint: 49% respond to only one. Flip volume, affected units and net upgrade value must be measured separately.

C.1 Common states: extra value of the BF16 action

Common-state comparison: details. Above, each design selects its own divergent states. A pre-registered comparison removes that difference: on 510 ordinary q8 states (85 tasks × 6 blocks, sampled by a fixed rule), the current action comes from B, from Q (GPTQ-W4, first checkpoint) or from Q′ (second checkpoint), and B or Q continues. $K(c) = V(\text{B}, c) - V(\text{Q}', c)$ asks what B's action adds beyond another W4 action. A frozen rule claims value only if both intervals exclude zero (Holm), equivalence only if both lie within ± 3 pp (two one-sided tests). The design first ran three blocks, and that first look was unresolved (under Q's continuation $K = -0.4$ pp, 95%: t [-9.3, $+8.5$], bootstrap [-3.5, $+2.7$]; under B's $+2.4$ pp, [-3.5, $+8.2$], [-0.8, $+5.5$]). Three more blocks were then added. That decision followed the unresolved first look, so a dated amendment fixed the extension before any new outcome existed, and a two-look reading (Pocock levels from a normal approximation: about 97% intervals for a difference and 94% for equivalence), was designated confirmatory in that amendment. The Pocock constant is not calibrated for small-k t intervals, the task bootstrap or the joint rule, so the two-look reading carries no exact error guarantee; we therefore report it as a sensitivity analysis next to the ordinary intervals and claim neither an advantage nor equivalence unless both agree. After the fixed extension (six blocks), B's action is not shown to be worth more than Q′'s, and equivalence is not established. Under Q's continuation, $K = +1.0$ pp (95%: t [-3.5, $+5.4$], bootstrap [-1.2, $+3.1$]; two-look 97%: [-4.2, $+6.2$], [-1.4, $+3.5$]). Under B's, $K = +0.8$ pp (95%: [-1.8, $+3.3$] for both; two-look 97%: [-2.2, $+3.8$], [-2.0, $+3.7$]). For equivalence under B's continuation, the ordinary 90% intervals lie within ± 3 pp (t [-1.2, $+2.8$], bootstrap [-1.4, $+2.9$]) and the two-look 94% intervals of the sensitivity analysis do not (t [-1.6, $+3.2$]); because the extension followed the first look, we do not claim equivalence. Without the 26 states at which earlier studies had already produced outcomes, K is $+0.5$ and $+0.8$ pp, also not resolved. Under Q's continuation, B's action flips $M = 7.1\%$ of outcomes (rescue 4.1, damage 2.9) for a net $\Delta = +1.2$ pp (95%: t [-1.3, $+3.6$]); Q′'s flips 8.0% (4.1, 3.9) for $+0.2$ pp ([-2.3, $+2.7$]; within ± 3 pp by two one-sided tests). A second W4 checkpoint's action departs from Q's at least as often as B's does (agreement with Q's action 53.9% against 55.7%; B with Q′ 57.3%; Figure 2d). The interaction of K with the continuation is -0.2 pp (95%: t [-4.5, $+4.1$], bootstrap [-2.7, $+2.4$]; secondary, not resolved). *q4, a separate prospective analysis.* The same design on Qwen3-4B planned three blocks. The frozen quarantine rule removed f01 (below), leaving 170 states. $K = -1.2$ pp under Q's continuation (bootstrap 95% [-5.9, $+3.5$]) and -2.4 pp under B's ([-6.5, $+1.2$]), both not resolved. With two blocks the t interval has one degree of freedom. Under Q's continuation the two block means happen to coincide, so it collapses to a point. That is a coincidence of two values, not a precise measurement or zero variance, and the reading rests on the bootstrap. *Identity checks.* Across both

models, 1,238 of 1,240 reference and implicit identity checks passed. Both failures are one q4 state. There the substituted action was semantically identical to the original after parsing, but its token sequence differed in the order of keys inside a tool-call argument. The canonical comparison treats such actions as equal, but the chat template renders the order, so the next prompt differed and the continuation could not be a replay. Every check with a token-identical action passed (869/869). As the frozen rule requires, the block was removed. Changing outcomes is thus not improving them; this does not make precision irrelevant or the two actions equivalent.

D From Numerical Fidelity to Task Value

These tests ask whether numerical error locates changed actions, consequential changes and beneficial changes. Methods share ordinary BF16-sourced states (U2B); the separate Mistral population is specified in Appendix E.3.

Contrasts. Methods are compared on shared states by $D_{cc} = Y_{\mathrm{QB}} - Y_{\mathrm{BB}}$ ($= -\Delta_{\mathrm{B}}$): Q's action against B's under the common BF16 continuation. Because BB is the shared BF16 original, two methods differ on the same states only through Y_{QB}. One pre-registered Holm family holds three contrasts on D_{cc}, with t intervals over block-mean differences: P1, QuaRot-W4 − GPTQ-W4; P2, SpinQuant-W4A4 − QuaRot-W4A4; P3, KIVI-2 − KV8. Unless marked crossed, intervals for these panels in §6 and §7 are fixed-pool 95% (t over block means, as fewer than 10 blocks are complete); other new-task (crossed) intervals are in the analysis outputs.

Numerical quality. QuaRot-W4 against GPTQ-W4 is -0.4 [-10.7, $+9.9$] pp (P1). The matched SpinQuant-W4A4 contrast is P2 (complete family: Table 10; paired diagnostic components: Table 9). Its third block rests on a one-time re-run of arms lost to a user-simulator overload; on the two blocks without it the contrast is $+2.4$ pp (strict view).

Upgrade modes on the q8 panel. On the q8 ordinary-state panel, B's action alone (G) is resolved in no cell. Fixed-pool intervals for B's takeover H or full substitution F exclude zero in weight-only (AWQ-W4, QuaRot-W4), W4A4 and KV8 cells, and new-task intervals still do in both W4A4 cells and KV8 (unadjusted), so resolved value is not confined to quantized activations. It is largest in SpinQuant-W4A4 ($H = +11.8$ [6.7, 16.8], $F = +13.7$ [9.3, 18.2] pp).

D.1 Fidelity, sensitivity and signed value

The following decomposition was added after panel results were read and is exploratory. On each cell's admitted states, exact action equality and deterministic continuation give

$$M_{cc} = P(a_{\mathrm{Q}} \neq a_{\mathrm{B}})\, P(Y_{\mathrm{QB}} \neq Y_{\mathrm{BB}} \mid a_{\mathrm{Q}} \neq a_{\mathrm{B}}).$$

B's rescues $r_{cc} = P(Y_{\mathrm{BB}} = 1, Y_{\mathrm{QB}} = 0)$ and damages $d_{cc} = P(Y_{\mathrm{BB}} = 0, Y_{\mathrm{QB}} = 1)$ give $D_{cc} = d_{cc} - r_{cc}$. The decomposition separates how often an action changes from what a changed action does; it is diagnostic rather than a mechanism.

Approximate KL. Both models are teacher-forced on identical prompt ids and B's action tokens. At each token, $\mathrm{KL}(\mathrm{B}\|\mathrm{Q})$ uses the union of both top-20 sets and the forced token, plus one tail bucket per model. An in-union token missing from a model's top-20 receives min(the 20th probability, tail mass divided by the number of missing tokens). The state score averages over B's action tokens. This coarse distribution is not the full token distribution.

Across and within configurations. Across eleven q8 cells, KL's descriptive Spearman correlation is 0.99 with the action-difference rate, 0.925 with M_{cc}, 0.81 with rescues and 0.93 with damages. It is 0.46 with the flip rate conditional on a difference and -0.44 with D_{cc}. Removing four unquantized path controls gives 0.98, 0.64 and -0.52 for the action-difference rate, conditional flip rate and D_{cc}. These configurations are not exchangeable draws. Within cells, the mean AUC is 0.93 for an action difference, 0.46 for a flip conditional on a difference and 0.55 for a rescue conditional on a flip; the last uses 3–39 flips per cell. These diagnostics establish no reliable predictor of net upgrade value.

KIVI. KIVI (Liu et al., 2024) quantizes keys per channel and values per token. KIVI-2 against KV8 is not resolved (-0.6 [-38.0, $+36.8$] pp), nor is KIVI-4 against KV8 ($+1.6$ [-2.9, $+6.0$] pp). Our implementation is a replay-compatible variant: at every decode step it reconstructs the quantized view that official KIVI's token-by-token decode

SpinQuant − QuaRot	difference [fixed pool]	[new tasks]
KL ($\times 10^{-3}$)	-34 [-57, -11]	[-56, -11]
Action difference (pp)	$+0.8$ [-11.4, $+13.0$]	[-8.2, $+9.4$]
Flip given difference (pp)	$+2.7$ [-10.2, $+15.3$]	[-7.5, $+11.4$]
D_{cc} (pp)	$+1.2$ [-3.9, $+6.2$]	[-4.7, $+7.1$]

Table 9: Paired diagnostic on 255 identical W4A4 states; 95% intervals, t over three block differences and crossed task×block bootstrap. The KL reduction is resolved; the differences in the task quantities remain unresolved. The D_{cc} row is primary contrast P2.

Contrast	blocks	ΔD_{cc} [fixed pool]	[new tasks]	Holm p
P1: QuaRot − GPTQ W4	3	-0.4 [-10.7, $+9.9$]	[-5.9, $+5.5$]	1.0
P2: SpinQuant − QuaRot W4A4	3	$+1.2$ [-3.9, $+6.2$]	[-4.7, $+7.1$]	1.0
P3: KIVI-2 − KV8	2	-0.6 [-38.0, $+36.8$]	[-7.1, $+7.1$]	1.0

Table 10: Complete primary method-comparison family on shared states, pp; fixed-pool and new-task 95% intervals. All three comparisons remain unresolved. Admission and re-run dependencies are described in the text.

would hold at that position from the full-precision cache, rather than persisting a compressed cache across turns, so every token, prompt or generated, attends with that view. The KIVI-2 contrast (P3) rests on two blocks: the third lost one arm to a server timeout that recurred on a single-request re-run, and both admitted blocks depend on the re-run of arms lost to server failures (strict rule: not estimable). Net of the kernel path, the quantization effect is $+0.6$ [-51.7, $+52.9$] (KIVI-2) and $+2.4$ [-5.4, $+10.1$] pp (KIVI-4), and the unquantized kernel path itself has $D_{cc} = -1.2$ [-6.2, $+3.9$].

D.2 All admitted ordinary-state upgrade effects

Table 11 reports all three upgrade modes for every admitted configuration, with both fixed-pool and new-task intervals.

D.3 Error dose: action agreement and end-to-end success

Supplementary perturbation analyses (exploratory). These analyses perturb the quantized weights directly instead of forking trajectories. They are exploratory; the main text uses only the error-dose rows (§6). Each derived checkpoint takes every decoder linear from the BF16 model, from Q, or from an interpolation between them; all other tensors come from the BF16 model, and the endpoints reproduce BF16 and Q bit for bit. Only finished end-to-end runs enter: q4 and q8, 85 tasks on each of three nonce blocks per cell (q3–q5 for W3, the development blocks f01–f03 for W4), with outage episodes excluded and context overflow counted as failure. Intervals are fixed-pool t intervals over the three blocks, with crossed task×block intervals alongside; there is no correction for multiplicity.

Error dose. The residual intervention $W_\alpha = W + \alpha(W_\mathrm{Q} - W)$ scales the quantizer's own weight error, with $\alpha \in \{0, 0.25, 0.5, 1, 2\}$; it is not a bit width (Table 12). Agreement with BF16's actions fell at every step of α in all six cells, with both intervals of every step below zero; in q8 RTN-W3, the share of BF16's decisions before the first differing action is 1, 0.32, 0.16 and 0.04 at $\alpha = 0$, 0.25, 0.5 and 1. Success did not follow. No success step up to $\alpha = 0.5$ is resolved under both intervals (in RTN-W3 the point estimates at $\alpha = 0.25$ lie above BF16, 40.0 and 36.5% against 35.7 and 30.6%), and success fell between $\alpha = 0.5$ and 1 only in RTN-W3, by -20.8 [-22.5, -19.1] pp in q4 (crossed [-29.8, -11.8]) and -14.1 [-26.9, -1.4] in q8 (crossed [-24.7, -3.9]). In GPTQ-W3, q4 GPTQ-W4 and q8 AWQ-W4 no success step is resolved under both intervals; doubling the q8 AWQ-W4 error ($\alpha = 2$) changed success by -10.6 [-16.4, -4.7] pp (crossed [-20.8, 0.0]). The $\alpha = 2$ arms of the four W3 cells did not finish. The action response to the error dose is thus monotone, whereas success shows only a drop between half and full RTN-W3 error.

D.4 The frozen trigger and its headroom

Following Lotfi et al. (2026), we scored every Bank U f01 state at the token level (1140 cell × state rows, 58 flips), teacher-forcing both models on the generated ids of a_B and a_Q. We record truncated top-20 entropy, the top-1 margin and a top-20-union KL, not full entropy or KL. All analyses are exploratory, and f01 is development data.

Uncertainty is associated with divergence. Q's entropy at the start of its action separates states where B and Q act differently from states where they act identically, with AUC 0.66–0.91 per sub-bank and model. Over all ordinary states, flips also sit at lower Q margins (U1: AUC 0.31 [0.22, 0.40] for q4, 0.24 [0.13, 0.36] for q8; stratified). Among already divergent U1 states, no Q-visible statistic has an interval excluding AUC 0.5; small-subset associations vary across state sources. This development diagnosis motivated the frozen margin rule, whose validation reuses the tasks with fresh

Cell	blocks	states	G	H	F
AWQ-W4	3	255	$+2.7\,[-4.6, +10.1]$ $[-2.4, +9.0]$	$+4.3\,[-0.1, +8.8]$ $[-1.2, +11.0]$	$+5.5\,[+1.0, +10.0]$ $[-0.8, +12.5]$
GPTQ-W4	3	255	$+0.4\,[-7.0, +7.7]$ $[-4.3, +5.9]$	$+0.8\,[-5.3, +6.9]$ $[-5.9, +7.5]$	$+3.5\,[-2.3, +9.4]$ $[-3.9, +10.6]$
KIVI-16 path	3	255	$+0.0\,[-2.9, +2.9]$ $[-2.7, +2.7]$	$+1.2\,[-6.6, +8.9]$ $[-3.5, +5.9]$	$+2.4\,[-0.6, +5.3]$ $[-2.4, +7.1]$
KIVI-2	2	170	$-2.9\,[-25.4, +19.5]$ $[-8.2, +2.4]$	$+2.9\,[-4.5, +10.4]$ $[-3.5, +10.0]$	$+2.9\,[-34.4, +40.3]$ $[-5.9, +10.6]$
KIVI-4	3	255	$+0.0\,[-5.1, +5.1]$ $[-3.5, +4.3]$	$+2.0\,[-1.4, +5.3]$ $[-3.5, +7.5]$	$+0.8\,[-8.1, +9.7]$ $[-5.9, +7.1]$
KV8	3	255	$+1.6\,[-0.1, +3.3]$ $[-1.2, +4.7]$	$+5.1\,[+0.6, +9.6]$ $[+0.4, +9.8]$	$+5.5\,[+1.0, +10.0]$ $[+0.0, +10.6]$
QuaRot-W16 path	3	255	$+0.0\,[-5.1, +5.1]$ $[-3.5, +3.1]$	$+1.6\,[-14.5, +17.7]$ $[-5.9, +9.4]$	$+2.4\,[-8.2, +12.9]$ $[-4.3, +9.0]$
QuaRot-W4	3	255	$+0.8\,[-5.3, +6.9]$ $[-3.5, +5.9]$	$+1.6\,[-1.8, +4.9]$ $[-4.3, +7.5]$	$+4.7\,[+1.8, +7.6]$ $[-2.0, +11.4]$
QuaRot-W4A4	3	255	$+1.2\,[-6.6, +8.9]$ $[-4.7, +6.3]$	$+6.3\,[-8.4, +21.0]$ $[-2.4, +14.1]$	$+9.4\,[+1.7, +17.1]$ $[+2.0, +16.5]$
SpinQuant-W16 path	3	255	$-0.4\,[-2.1, +1.3]$ $[-3.1, +2.4]$	$+2.0\,[-1.4, +5.3]$ $[-2.4, +6.3]$	$+2.4\,[-0.6, +5.3]$ $[-2.4, +7.1]$
SpinQuant-W4A4	3	255	$+3.1\,[-2.9, +9.2]$ $[-2.0, +8.2]$	$+11.8\,[+6.7, +16.8]$ $[+5.1, +18.4]$	$+13.7\,[+9.3, +18.2]$ $[+5.5, +21.6]$

Table 11: All q8 ordinary BF16-sourced states admitted by the panel rules, including unquantized path controls. Effects in pp; each cell: estimate with fixed-pool 95% t interval, then crossed task×block 95% interval. All three modes use the same admitted states within a row. These descriptive intervals are unadjusted; the primary method family is P1–P3. KIVI-2 admits only two blocks after one-time infrastructure re-runs; it is not estimable under the original strict rule. Other cells admit three blocks. No action-only interval excludes zero; several takeover or joint-upgrade intervals do.

Cell	$\alpha = 0$ (BF16)	0.25	0.5	1 (Q)	2
q4 RTN-W3	35.7 (1)	40.0 (0.362)	33.7 (0.225)	12.9 (0.088)	n.f.
q8 RTN-W3	30.6 (1)	36.5 (0.319)	27.5 (0.162)	13.3 (0.043)	n.f.
q4 GPTQ-W3	35.7 (1)	37.6 (0.508)	36.1 (0.406)	36.1 (0.234)	n.f.
q8 GPTQ-W3	30.6 (1)	29.4 (0.397)	29.0 (0.297)	26.7 (0.176)	n.f.
q4 GPTQ-W4	33.7 (1)	38.0 (0.572)	38.4 (0.504)	43.9 (0.424)	37.6 (0.238)
q8 AWQ-W4	31.4 (1)	31.0 (0.392)	32.5 (0.309)	32.5 (0.194)	22.0 (0.058)

Table 12: Error dose (exploratory): closed-loop success (%) of W_α and, in parentheses, action agreement with BF16 (share of BF16's decisions before the first differing action; 1 = identical). 255 task×block units per cell over three blocks; action agreement on 245 units for q4 RTN-W3, whose other 10 units contain an episode that ended in context overflow and has no message record. n.f.: not finished. Intervals of every step are in the analysis output.

nonce blocks.

D.5 Matched references for the frozen trigger

Table 13 compares the frozen trigger of §6 (frozen rule, unchanged) with two references on the same 1020 f02/f03 U1 rows, which hold 38 rescues and 31 damages. The *matched action-replacement oracle* selects, per cell, exactly as many rows as the frozen rule did, taking only rows with $G > 0$ and filling with $G = 0$ rows when a cell has too few rescues; it never selects a damage. It is hindsight selection, not a policy. At the 10% and 20% budgets every rescue fits. The *three-option menu oracle* $E[\max(G, H, F, 0)]$ is the share of states where some arm beats QQ. On Bank P divergent states it is 0.202 over the four W4 cells (state-weighted, the one exception to the paper's equal-weight pools) against 0.117 for the action alone (1127 vs. 654 of 5576 states), and 0.191 vs. 0.113 over the six core cells. It is hindsight selection that also allows different replacement types, with different costs, so it is not a same-cost deployment comparison and not comparable with the action rule.

The *KL reference row* is a post-hoc auxiliary baseline. Its score is a truncated $\mathrm{KL}(\mathrm{B}\|\mathrm{Q})$ over the union of both models' top-20 next-token sets plus a tail bucket, per a_Q token, teacher-forced on identical prompt ids. It needs a BF16 prefill of the prompt and a_Q at every decision (here 6709.3 prompt and 82.3 action tokens on average), so it is not Q-only. It captures turnover above chance (33 vs. 16.7 at 20%) but no significant net gain (Holm $p = 1.0$). Three aggregations over the a_Q tokens were fit on f01 only, with per-cell thresholds and the frozen rule's selection criterion; the maximum was chosen, and all of f02/f03 is reported. Because the f02/f03 outcomes had been seen, the row is post hoc, and it is not a DMR reproduction, which trains a router.

Budget	same per-cell selection counts sel.	oracle r/d/net (pp)	frozen r/d/net (pp)	captured	KL reference (max. per-token KL; needs BF16; post hoc) sel.	net vs. random (p; Holm)	r/d	turnover vs. random (p)
5%	50	34/0/+34 (+3.33)	6/2/+4 (+0.39)	4/34	72	−1 vs. +0.45 (.8337; 1.0000)	5/6	11 vs. 4.58 (.0038)
10%	114	38/0/+38 (+3.73)	7/5/+2 (+0.20)	2/38	126	+1 vs. +0.82 (.5483; 1.0000)	8/7	15 vs. 7.54 (.0048)
20%	222	38/0/+38 (+3.73)	16/11/+5 (+0.49)	5/38	257	+3 vs. +1.59 (.3984; 1.0000)	18/15	33 vs. 16.73 ($< 10^{-4}$)

Table 13: Frozen Q-only trigger against a matched action-replacement oracle (hindsight; same per-cell selection counts) and a post-hoc KL reference row, on f02/f03 U1 (1020 rows). Captured: frozen net over oracle net. Random: matched-count per-cell random-trigger expectation; p: exact one-sided hypergeometric, Holm over the three budgets.

E Where Do the Consequences Persist?

E.1 Ordinary states and the user simulator

Ordinary states turn over less, mostly through text under our simulator. At ordinary states, replacing Q's action changes the outcome less often than at first tool-structure divergences (a descriptive comparison of different state populations, with no interval on the difference). Over blocks f01–f03 (task-equal; t intervals over block means), M per sub-bank is 7.8–9.0% for the q4 W4 cells in U1 and U2B, 4.7–9.0% for the q8 W4 cells and at most 4.7% for KV8. The net gain lies within ±2.8 pp (U1, U2B) and ±5.0 pp (U2Q), and every interval covers zero (Appendix E.2). Identical actions never change the outcome, as determinism implies. Reused arms rest on canonical equality. All 4,114 such reuses were rendered with each model's chat template; 4,113 are token-identical. The remaining pair differs only in JSON key order, and its two cross arms, measured separately, gave the reused outcomes. After deduplication, 150 of the 208 outcome-changing cell×state rows (≈72.1%) are class (ii), the same tool structure with different text; 26 involve a different tool name, call count or argument (Table 15). The rows are not independent states: the three blocks repeat the same tasks, and each U2B/U2Q state appears once per cell of a model. This is a *composition share*: per instance, text, tool/argument and act/talk differences flip the outcome in 14.4%, 18.2% and 18.9% of rows. Text dominates because it is common (1045 of 1357 differing rows), not because each text difference matters more.

Sensitivity to the user simulator. Text actions act through the user simulator, and changing the simulator moves τ-bench outcomes by 9–15 points (Seshadri et al., 2026; Jain and Sandhu, 2026). We froze a control before running it: the 39 deduplicated class-(ii) flips of block f01 and 39 matched non-flip states. Their BQ and QQ branches were each re-run three times with a different user simulator, Sonnet 5.5, answering through a file relay from Claude Code subagents (reproducibility limits in the Limitations). Judged by the sign of the replicate mean, only 7 of the 39 flips keep their direction (95% CI 0.075–0.335). Six reverse, and 26 vanish (15 both succeed, 2 both fail, 9 ties). Among the matched non-flip states, 11 of 39 become flips (0.15–0.45). Replicates of the same arm disagree in 41% of flip arms and in 21–26% of control arms. **The Bank U text-difference outcome labels did not reproduce reliably under this alternative relay setting.** The control changes three things at once (simulator model, sampling, relay interface), so it cannot say which of them changed the labels.

Sensitivity of the four-branch results. Before running, we sampled 80 Bank P states without outcome information (20 per W4 cell, stratified by domain and block) and re-ran all four arms five times each with Sonnet 5.5 as the user, one isolated single-request call per user turn (7896 calls; all integrity checks passed). Per-state arm outcomes track the original single runs (correlation 0.63–0.68). Pooled over cells, $G = -3.0$ $[-10.2, +4.8]$, $H = +5.5$ $[+0.0, +11.2]$ and $G - H = -8.5$ $[-14.8, -2.8]$ pp under Sonnet, against $G - H = -2.5$ $[-11.2, +6.2]$ on the same states under the original simulator (change -6.0 $[-16.2, +4.2]$). With 20 states per cell, this control cannot resolve the few-pp contrasts of Table 5 (on this sample even the original simulator gives $G - H = 0.0$ in both q8 W4 cells), so it neither confirms nor reverses the within-cell orderings.

Bank P and Bank U measure different quantities. Bank P skips text-only differences by design, yet these make up 150 of the 208 outcome-changing rows at ordinary states, and their labels did not reproduce under a second simulator (Appendix E.2; §7). Including text actions requires a simulator control before their consequences are reported as a property of the action.

E.2 Bank U per Sub-Bank

Table 14 gives the per-state turnover behind Appendix E.1, over blocks f01–f03. Exact action agreement is 0.43–0.83 per cell and sub-bank.

Cell	U1 (n=255)	U2B (n=255)	U2Q (n=60)
q4 GPTQ-W4	.090 (11/12)	.082 (10/11)	.083 (4/1)*
q4 AWQ-W4	.082 (7/14)	.078 (7/13)	.083 (3/2)
q4 KV8	.047 (4/8)	.039 (5/5)	.017 (0/1)
q8 GPTQ-W4	.047 (8/4)	.067 (9/8)	.050 (3/0)*
q8 AWQ-W4	.086 (14/8)	.090 (15/8)	.067 (3/1)
q8 KV8	.035 (5/4)	.031 (6/2)	0 (0/0)

Table 14: Bank U, blocks f01–f03: per-state turnover M (rescues/damages), task-equal (one state per task and block; 85/85/20 states per block). G lies within ± 0.028 in U1 and U2B and within ± 0.050 in U2Q, and every interval includes 0. Intervals are t intervals over the three block means and are wide (e.g. q4 GPTQ-W4 U1: .090 [.029, .151]). *GPTQ-W4's U2Q states are its own U1 states, with concordant outcomes on 20/20 duplicates per block; they are counted once in cross-sub-bank totals.

Model:sub	(i) identical	(ii) same tool structure, text	(iii) different tool or args	(iv) act vs. talk
q4:U1	399: 0	309: 45 (.146)	19: 1 (.053)	38: 10 (.263)
q8:U1	483: 0	202: 28 (.139)	47: 11 (.234)	33: 4 (.121)
q4:U2B	430: 0	277: 44 (.159)	24: 4 (.167)	34: 3 (.088)
q8:U2B	470: 0	199: 25 (.126)	40: 9 (.225)	56: 14 (.250)
q4:U2Q	108: 0	63: 10 (.159)*	5: 0	4: 1 (.250)
q8:U2Q	129: 0	33: 6 (.182)*	11: 1 (.091)	7: 0
unique rows	1943: 0	1045: 150 (.144)	143: 26 (.182)	169: 32 (.189)

Table 15: Which kind of action difference changes the outcome at ordinary states (Bank U, blocks f01–f03). Entries: "state rows: outcome-changing rows (flip rate)", pooled over the three cells of a model within a sub-bank. The last row counts each (cell, state) once: q4/q8 GPTQ-W4's U2Q states repeat their U1 states (* includes 5 and 3 repeated class-(ii) flips), so 216 flipping rows become 208. In block f01 alone the corresponding counts were 39 of 56 flips in class (ii), with rates .126/.109/.200 for (ii)/(iii)/(iv); the 58 f01 flips scored in Appendix D.4 include the two duplicated U2Q rows.

E.3 Models compared on their own ordinary states

Mistral population, scoring and coverage. The Mistral panel applies the same row rule to Mistral's own BF16 trajectories on the 85 retail and airline tasks. Seven BF16 source rows are missing because the source episode ended in an invalid tool call produced by the model itself (airline task 44 in all six blocks, airline task 13 in f01). A dated amendment, fixed from the BF16 source runs before any Mistral fork output existed, excludes both tasks in every block and cell, leaving 83 tasks. The selection was outcome-blind, but BF16 Mistral fails these tasks structurally, so the estimates describe this conditional population, not the 85-task panel, and are not compared with q8 as if the populations were equal. A branch that ended because the harness could not accept the agent's own message (an empty or invalid tool name, or an empty message) is scored as a failure, as in end-to-end evaluation. This rule was added after the Mistral phase-1 results had been read; without it the Mistral primary comparison is not estimable from phase 1 (one complete block). The analysis without the rule (which leaves such an arm missing and fails its block) and one applying the rule only to data not yet read when it was fixed are reported alongside; both admit the same three blocks for GPTQ − RTN ($+1.2$ [-13.8, $+16.2$]). Coverage: GPTQ-W4 has 498 states on six blocks and RTN-W4 415 states on five. Full estimates on the 83 tasks (GPTQ-W4 on six blocks, RTN-W4 on five; Table 16): in GPTQ-W4, B's action rescues $r = 4.6$ [2.4, 6.8] and damages $d = 5.4$ [2.4, 8.5] pp of states ($M = 10.0$ [7.0, 13.1]; fixed-pool 95%), and no mode's net value is resolved ($G = -0.8$ [-5.2, $+3.6$], $H = +1.4$ [-3.0, $+5.8$], $F = +0.0$ [-4.2, $+4.2$]). RTN-W4 is similar ($M = 10.4$ [6.3, 14.4]). The preregistered contrast GPTQ − RTN is not resolved: $\Delta D_{cc} = +2.2$ [-3.5, $+7.8$] pp, $p = .35$, on five blocks, two of which are admitted only by a scoring rule adopted after an interim look (above).

Cell	blocks	states	G	H	F
GPTQ-W4	6	498	-0.8 [-5.2, $+3.6$] [-5.8, $+4.2$]	$+1.4$ [-3.0, $+5.8$] [-4.0, $+6.8$]	$+0.0$ [-4.2, $+4.2$] [-6.0, $+6.0$]
RTN-W4	5	415	$+0.7$ [-2.7, $+4.2$] [-3.6, $+5.5$]	$+1.9$ [-3.8, $+7.7$] [-4.6, $+8.0$]	$+2.4$ [-2.2, $+7.0$] [-3.9, $+8.4$]

Table 16: Mistral-Small-3.2-24B on its fixed 83-task conditional population. pp; fixed-pool 95% t intervals, then crossed task×block 95% intervals. Missing argument-JSON errors leave RTN f04 excluded, whereas model-produced invalid tool names are scored as failures under the dated amendment above. None of the three mode intervals excludes zero.

E.4 Replacement and four-branch estimands on BFCL and ALFWorld (supplementary)

Table 17 gives the replacement quantities of §4 (r, d, M and G_{all} over all items, non-divergent items counting as zero) and applies the four-branch analysis to the external benchmarks, from the raw

four-arm outcomes on divergent items (one run per item; item bootstrap within benchmark category). The four-branch columns are post hoc and not part of any pre-registered analysis. In the KV4 cells $Y_{\mathrm{QQ}} = 0$ on every state, so Δ_{Q} is almost constant and their off-diagonal mass is H. In BFCL q4 GPTQ-W3, the base category was dropped whole by the QQ no-op rule, leaving 591 items and 540 states. Off-diagonal shares are 4.7–17.1% in BFCL W4 cells (pool 10.3%) and 32.8–40.0% in q8 ALFWorld W4 (pool 36.4%). The ALFWorld GPTQ-W3 stress cell has $H = +38.5$ and $G - H = -34.6$ [-43.5, -25.4] pp; KV4 cells reduce to H.

E.5 Telecom development and task concentration

On 80 telecom development tasks (6 blocks; fixed-pool 95%), the two q8 holdout cells give $G_{\mathrm{all}} = +2.7$ [-4.7, $+10.1$] (GPTQ-W4) and -0.6 [-4.9, $+3.7$] pp (AWQ-W4), against -5.6 and -6.6 on the holdout; q4 has no telecom development forks.

On the 85 development tasks, the top 20% of tasks account for 57–71% of flips in the W4 and KV8 cells. Splitting the 20 blocks into disjoint nonce halves reproduces task-level turnover rankings (Spearman 0.52–0.86). Conditional tests control measured baseline success and, in a post-hoc variant, trajectory length; the association is carried primarily by retail and weakens on airline alone. These are retrospective associations on the same tasks under fresh prompt perturbations, rather than evidence of transfer to new tasks.

Split-half fragility (Table 1). The halves are assigned by an outcome-independent hash of the block name (sha256 of a fixed prefix and the block name, mod 2), written to a file before this post hoc analysis loaded outcomes; fragility is a task's mean turnover over the four W4 cells and the defining half's blocks, and ties at the top-20% threshold are broken by a hash. In the reverse direction (rank on the second half, estimate on the first), the 17 fragile tasks turn over 39.0% (t [34.2, 43.9]) and carry 42.3% of the held-out turnover, the other tasks 13.3%; the fragile tasks' net gain is $+3.7$ pp (t [$+0.6$, $+6.9$]; crossed [-5.0, $+12.4$]) and the rest's $+0.5$. Ranking tasks by BF16 success instead (tertiles on one half) also separates turnover: the hardest tertile turns over 6.7–10.0% and the others 25.4–27.3%, with no net gain resolved under the crossed interval.

E.6 Takeover duration

Design and reading. Protocol `TAKEOVER_WINDOW_PROTOCOL.md` (tw-v1, exploratory; analysis code frozen before outcomes were read, Amendments 2–4). At 255 q8 BF16-sourced states (blocks f01–f03, 85 tasks each), the current action comes from Q (or B); BF16 executes the next L complete decisions and hands back to Q. $V_{Q,L} = Y_{Q,L} - Y_{Q,0}$ is the value of an L-decision BF16 window after Q's action, $R_{Q,L} = Y_{Q,L} - Y_{Q,\infty}$ its shortfall against sustained takeover. All six tasks passed acceptance (0 errors, 0 overflows, 256/256 endpoint re-runs reproduced, nothing quarantined). The $E_{Q,4}{=}1$ subset keeps states where continuous BF16 generated more than four decisions, so a four-decision window really hands back. Intervals: block-t over 3 block means (registered primary, df $= 2$) and crossed task$\times$block bootstrap; we read a result as resolved only when both exclude zero (Table 18). The four-decision gain on the $E_{Q,4}{=}1$ subset is resolved under block-t but not under the crossed interval. Differences between action sources at each L (Δ_L) are unresolved in both cells.

F Independent Telecom Confirmation

The protocol was frozen before any holdout unit was run, and its hash was recorded in the split manifest. It fixes the following design:

- tasks, cells, weights and nonce blocks;
- two fixed halves of six blocks each;
- an outcome-blind repair rule, with reserve blocks used only for infrastructure repair;
- the decisional crossed interval (two-way ANOVA moments, normal 90%);
- the readings: positive if $L > 0$, above 3 pp if $L > 3$, small if the interval lies inside ± 3 pp, otherwise inconclusive;
- the secondary families S1 (both $r > 0$ and $d > 0$) and S2 (persona-stratified permutation, 10,000 draws, no outcome screen);
- six stratum families, with Holm correction within each.

The analysis script refuses holdout data unless the hash matches. Two task-cell pairs were excluded as structural, and the complete-case analysis is identical.

	all items					divergent items					
Cell	items	r	d	M	G_{all}	states	G	H	F	$G-H$	off-diag. %
BFCL q4 GPTQ-W4	788	6.3	4.3	10.7	+2.0 [−0.3, +4.3]	573	+2.8 [−0.4, +5.9]	+1.4 [−0.7, +3.5]	+4.2 [+0.9, +7.2]	+1.4 [−1.8, +4.5]	11.7
BFCL q4 AWQ-W4	788	5.7	9.5	15.2	−3.8 [−6.5, −1.3]	642	−4.7 [−7.8, −1.5]	−1.1 [−3.6, +1.4]	−1.4 [−4.8, +1.7]	−3.6 [−6.7, −0.2]	17.1
BFCL q8 GPTQ-W4	788	2.9	2.3	5.2	+0.6 [−0.9, +2.3]	533	+0.9 [−1.3, +3.3]	+1.5 [+0.0, +3.2]	+2.1 [−0.4, +4.6]	−0.6 [−3.1, +2.0]	4.7
BFCL q8 AWQ-W4	788	4.1	2.7	6.7	+1.4 [−0.3, +3.3]	548	+2.0 [−0.4, +4.7]	+1.8 [+0.0, +3.7]	+3.6 [+0.9, +6.5]	+0.2 [−2.6, +3.1]	7.8
BFCL W4 pool					+0.1 [−1.0, +1.2]		+0.3 [−1.2, +1.8]	+0.9 [−0.2, +2.0]	+2.1 [+0.3, +3.9]	−0.6 [−2.3, +1.0]	10.3
ALFWorld q8 GPTQ-W4	134	9.0	12.7	21.6	−3.7 [−11.2, +4.5]	115	−4.3 [−13.3, +5.0]	+7.0 [−2.6, +16.7]	+9.6 [−0.9, +19.7]	−11.3 [−20.9, −1.8]	40.0
ALFWorld q8 AWQ-W4	134	10.4	14.2	24.6	−3.7 [−11.9, +4.5]	119	−4.2 [−13.6, +4.9]	+5.0 [−2.6, +12.2]	+4.2 [−5.9, +13.7]	−9.2 [−18.8, +0.0]	32.8
ALFWorld W4 pool					−3.7 [−9.7, +1.9]		−4.3 [−11.1, +2.2]	+6.0 [−0.5, +12.6]	+6.9 [−2.1, +15.5]	−10.3 [−17.7, −2.9]	36.4
BFCL q4 GPTQ-W3^s	591	6.1	3.6	9.6	+2.5 [+0.2, +5.1]	540	+2.8 [+0.2, +5.5]	+5.2 [+2.9, +7.7]	+8.9 [+5.7, +12.2]	−2.4 [−5.7, +0.7]	13.0
BFCL q8 GPTQ-W3^s	788	5.2	4.1	9.3	+1.1 [−1.0, +3.3]	696	+1.3 [−1.1, +3.7]	+1.4 [−0.1, +2.9]	+3.4 [+0.9, +6.0]	−0.1 [−2.6, +2.4]	8.3
BFCL q4 KV4^s	788	0.3	0.0	0.3	+0.3 [+0.0, +0.6]	787	+0.3 [+0.0, +0.6]	+2.9 [+1.9, +4.1]	+20.3 [+17.6, +23.0]	−2.7 [−3.9, −1.5]	20.5
BFCL q8 KV4^s	788	0.3	0.0	0.3	+0.3 [+0.0, +0.6]	788	+0.3 [+0.0, +0.6]	+1.8 [+1.0, +2.7]	+14.5 [+12.1, +17.0]	−1.5 [−2.5, −0.6]	14.7
ALFWorld q8 GPTQ-W3^s	134	6.7	3.0	9.7	+3.7 [−1.5, +9.7]	130	+3.8 [−1.5, +9.8]	+38.5 [+29.8, +46.9]	+42.3 [+33.6, +50.8]	−34.6 [−43.5, −25.4]	30.8
ALFWorld q8 KV4^s	134	0.0	0.0	0.0	+0.0 [+0.0, +0.0]	134	+0.0 [+0.0, +0.0]	+38.1 [+30.6, +45.5]	+53.0 [+44.8, +61.2]	−38.1 [−45.5, −30.6]	26.9

Table 17: Replacement and four-branch estimands on BFCL multi-turn and ALFWorld (no LLM user simulator; one run per item; 95% item-bootstrap intervals). All items: rescue r, damage d, turnover M (% of all items) and net G_{all} (pp), non-divergent items counting as zero. Divergent items (supplementary, exploratory, post hoc; pp, divergent items): Off-diag.: share with $\Delta_{\text{Q}} \neq \Delta_{\text{B}}$. Pools: equal weight per cell. ^{s}Stress and unrotated KV4 cells; in KV4 $Y_{\text{QQ}} = 0$ on every state, so the table reduces to H.

cell	states	$V_{Q,1}$	$V_{Q,4}$	$V_{Q,\infty}$	$R_{Q,4}$
GPTQ-W4	all 255	+2.4 [-5.4, +10.1; -3.1, +7.8]	+2.0 [-5.4, +9.3; -3.9, +8.2]	+0.8 [-5.3, +6.9; -5.9, +7.5]	+1.2 [-6.6, +8.9; -3.1, +6.3]
GPTQ-W4	$E_{Q,4}$=1	+3.4 [-8.9, +15.7; -4.7, +10.9]	+2.0 [-7.0, +11.0; -6.3, +10.6]	+0.2 [-8.2, +8.6; -9.9, +9.5]	+1.8 [-10.6, +14.1; -5.1, +9.6]
SpinQuant-W4A4	all 255	0.0 [-2.9, +2.9; -3.5, +3.5]	+6.3 [+0.2, +12.4; 0.0, +12.2]	+11.8 [+6.7, +16.8; +5.1, +18.4]	−5.5 [-7.2, -3.8; -10.6, -0.8]
SpinQuant-W4A4	$E_{Q,4}$=1	−1.8 [-6.2, +2.7; -5.5, +1.6]	+4.2 [+1.2, +7.2; -3.0, +11.6]	+12.5 [+7.2, +17.7; +4.6, +21.0]	−8.3 [-10.9, -5.6; -16.0, -1.2]

Table 18: Takeover duration (pp; [block-t; crossed] 95%). SpinQuant-W4A4 four-decision windows handed back at 66% of states (169/255) and generated about 228 BF16 tokens per state, against about 60 for one-decision windows.

Replay acceptance on the holdout (post hoc audit). Y_{QQ} is the one arm taken from the original episode rather than re-executed, and an exploratory count showed only-QQ successes more often than random labelling would give in two q8 holdout cells. We therefore audited the holdout's acceptance records after the analysis, with a script that outputs only consistency checks and counts. The fixed hash rule selected 135 of the 1202 divergent states for the QQ no-op; all 135 were run and passed, none was missing, failed or non-computable, and no other state was re-run. Beyond the frozen checks (suffix, prompts, termination, outcome, budget counters), the 135 also match on reward, database state, the official reward, the final environment hash against a strict replay of the original episode, and the history prefix. The original QQ and the re-run arms are scored by the same official reward rule: recomputing it from the stored reward information reproduces the stored value for every arm. The audit read the same input files as the frozen analysis, so no repair under the protocol's missing-data rules was needed, and the pre-registered result stands as reported. The audit does not explain the only-QQ pattern; a real upgrade that damages the original policy would also produce it. Re-including the two structurally excluded task–cell pairs in all twelve blocks (post hoc sensitivity, frozen estimator) with every unit set to -1 or $+1$ gives point estimates of G_{all} from -6.68 $[-10.35, -3.01]$ to -2.27 $[-5.96, +1.43]$ pp (crossed 90%). This is a range of point estimates: the interval of the most optimistic scenario includes zero. With their observed units kept in the eleven blocks where they exist and only the two missing units set to -1 or $+1$, the estimates are -4.72 to -4.35 pp, each with a crossed 90% interval below zero. The frozen reading is inconclusive in all 36 combinations. Table 19 gives the confirmatory result. Table 20 gives all six stratum families (Holm within family), which are descriptive. Every level estimate is ≤ 0, and the Holm-significant levels are negative: step-budget T2 and T3, mobile-data issues (-9.09 pp) and the fewest-faults tercile (-6.37 pp).

Cell (weight)	units	excl.	G	r	d	ρ_c
q4 GPTQ-W4 (1/2)	396	1	−3.03 [−7.86, +1.80]	+5.05	+8.08	0.668
q8 GPTQ-W4 (1/4)	408	0	−5.64 [−9.33, −1.94]	+4.41	+10.05	0.374
q8 AWQ-W4 (1/4)	396	1	−6.57 [−11.98, −1.15]	+6.31	+12.88	0.644
Weighted (G_{all})			−4.57 [−7.17, −1.96]	+5.21	+9.77	0.588

Table 19: Telecom holdout, pre-registered analysis (34 tasks × 12 blocks; pp; crossed 90% intervals). Per-cell rows are descriptive; the weighted row is decisional (G_{all}, reading: inconclusive) and S1/S2 (both pass). Excl.: structurally excluded tasks (decision/trajectory count mismatch at search).

Family	Level	n	G [crossed 90%]	p	Holm p
step budget[a]	T1	138/129/110	-0.32 [-1.57, $+0.94$]	.6771	.6771
	T2	124/126/162	-2.30 [-3.88, -0.72]	.0165	.0495
	T3	123/144/124	-1.94 [-3.35, -0.54]	.0229	.0495
persona	Hard	11	-5.99 [-12.23, $+0.26$]	.1147	.3355
	None	12	-3.65 [-7.42, $+0.13$]	.1118	.3355
	Easy	11	-4.36 [-9.19, $+0.48$]	.1385	.3355
BF16 difficulty	0	20	-2.35 [-4.09, -0.62]	.0257	.0772
	low	11	-4.54 [-8.90, -0.19]	.0860	.1720
	high	3	-18.75 [-38.00, $+0.50$]	.1091	.1720
task length	T1	9	-5.09 [-9.81, -0.38]	.0755	.1511
	T2	8	-0.78 [-4.46, $+2.90$]	.7270	.7270
	T3	17	-6.17 [-10.49, -1.84]	.0189	.0567
issue type	mms	14	-1.15 [-3.33, $+1.04$]	.3880	.4434
	mobile data	11	-9.09 [-14.98, -3.20]	.0111	.0333
	service	9	-4.40 [-10.32, $+1.52$]	.2217	.4434
faults	T1	19	-6.37 [-9.91, -2.83]	.0030	.0091
	T2	10	-1.94 [-7.58, $+3.69$]	.5704	.5704
	T3	5	-2.92 [-5.18, -0.66]	.0337	.0674

Table 20: Telecom holdout strata: the six pre-registered families (descriptive; Holm within each family over its three levels). pp. For task-level families, n is the number of tasks and G is G_{all} over the tasks in the level. [a]Step budget (remaining steps at the divergence) is unit-level: n is the number of divergent units per cell (q4 GPTQ-W4/q8 GPTQ-W4/q8 AWQ-W4), and G is the level's contribution to G_{all}, so the three levels sum to G_{all}. Tertile cut points were fixed on telecom development tasks.

Post-hoc four-branch estimands on the holdout. Table 21 is a supplementary, descriptive output computed after the confirmatory analysis, with the same admission and units. It is not part of the pre-registered confirmation and changes neither the primary endpoint nor its reading. Its G is a ratio over divergent states, whereas G_{all} in Table 19 averages over all admitted units with non-divergent units counted as 0.

G Boundaries of a Local-Defect Explanation

Controlled injection. On 30 ordinary GPTQ-W4 states per model, each from one development block, we compared a forced invalid action with a matched valid action and with a process perturbation that repeated the last successful read; either B or Q then continued. Each panel balanced wrong-ID, tool-call/schema error and premature-write perturbations. Table 22 gives the invalid-minus-process contrast. These small panels did not resolve an error-specific effect; the q8 interval under Q still allows substantial harm.

Replay checks passed on both panels; q4 lacked three recorded audit fields. In q4, most wrong-tool injections actually produced schema errors (8/10). In q8, seven premature writes failed to inject the intended write (six environment rejections and one execution without a database change), and two valid references returned "not found". These failures prevent the panels from establishing a general recovery mechanism.

G.1 Claims not supported by their controls

W×KV interaction (not supported). A first q4 analysis found a W4×KV4 interaction of -14.9 [-23.8, -6.0] pp ($k = 3$). Its KV16 reference lacked the Q/K rotation used for KV4, so rotation and quantization were confounded. With a matched rotated reference, the q4 interaction is -9.0 [-30.6, $+12.5$]. The pre-registered q8 replication gives $+2.1$ [-5.8, $+10.1$] over 6 blocks, crossed [-8.2, $+12.4$]. The replication is not supported, and we make no W×KV claim.

Static precision restoration. Restoring attention or a layer segment was compared with byte-matched random modules. In the damaged RTN-W3 models the attention contrast was $+9.8$ [$+6.1$, $+13.6$] pp (crossed [$+4.5$, $+15.1$]), while the last third was -8.6 [-13.4, -3.8] (crossed [-12.9, -4.7]). These were exploratory contrasts among 39 unadjusted comparisons; the attention reference had only two completed random draws. None of the 30 GPTQ-W3 or W4 contrasts excluded zero under both intervals, and no tested allocation beat whole-model GPTQ-W4 at matched storage. This incomplete stress test supports no deployment method. The full matrix and unfinished-arm inventory remain in the released analysis.

Low-bit stress failures. Failure changes include protocol and budget failures, not only a wrong final state. In q4 RTN-W3, 32.9% of finished episodes reached context or step limits, and 44.7% ended with required tool calls emitted as plain JSON that the frozen parser did not accept; none issued a parsed write. In q8 RTN-W3, these limit endings occurred in 20.4%, and 16.9% hit the failed-call cap. On the 24-task, one-block q8 W2 panel with a loop guard, 66.7–87.5% reached such limits. These descriptive observations depend on the parser, budget and stopping rules and do not identify a universal bit-width threshold.

Provenance. Figure values, retained tables and scoring definitions map to committed analysis outputs and frozen manifests. The release also retains complete pattern counts, restoration matrices, process labels, experiment logs and historical corrections; they are not additional independent evidence

Cell (weight)	states	G	H	F	$G - H$	mean I	off-diag. %
q4 GPTQ-W4 (1/2)	385	$-3.1\,[-9.0, +2.8]$	$+0.5\,[-5.4, +6.4]$	$+1.8\,[-4.8, +8.4]$	$-3.6\,[-9.1, +1.8]$	$-4.4\,[-12.5, +3.7]$	19.7
q8 GPTQ-W4 (1/4)	399	$-5.8\,[-10.2, -1.3]$	$-4.0\,[-10.2, +2.2]$	$-5.3\,[-11.6, +1.1]$	$-1.8\,[-6.4, +2.9]$	$-4.5\,[-10.4, +1.4]$	19.5
q8 AWQ-W4 (1/4)	396	$-6.6\,[-13.0, -0.1]$	$-8.3\,[-15.3, -1.3]$	$-8.3\,[-15.5, -1.2]$	$+1.8\,[-3.3, +6.8]$	$-6.6\,[-14.5, +1.4]$	23.7
Weighted pool	1180	$-4.6\,[-7.8, -1.5]$	$-2.8\,[-6.3, +0.6]$	$-2.5\,[-6.2, +1.3]$	$-1.8\,[-4.8, +1.1]$	$-5.0\,[-9.6, -0.4]$	20.7

Table 21: Telecom holdout, *post-hoc* four-branch estimands on divergent states (descriptive; not part of the pre-registered confirmation). pp; crossed task×block 95% intervals; pool weights as pre-registered.

Model	continuation	error minus process [95%]	failed checks
q4	B	$0.00\,[-0.13, +0.13]$	2/30
q4	Q	$-0.03\,[-0.20, +0.10]$	2/30
q8	B	$-0.03\,[-0.13, +0.07]$	9/30
q8	Q	$-0.13\,[-0.30, +0.03]$	9/30

Table 22: Error-specific contrast, in success probability units; task-cluster bootstrap. Failed manipulation checks are state counts, repeated here for the two continuations of the same panel. No equivalence claim follows from intervals including zero.

for the paper's conclusions. Protocol amendments and unresolved primary contrasts remain disclosed in this supplement. Removing duplicate displays and exploratory leads does not change the estimators, admission rules or reported primary decisions.